



Big Data, Techniques and Platforms

Cours 1/8: MapReduce

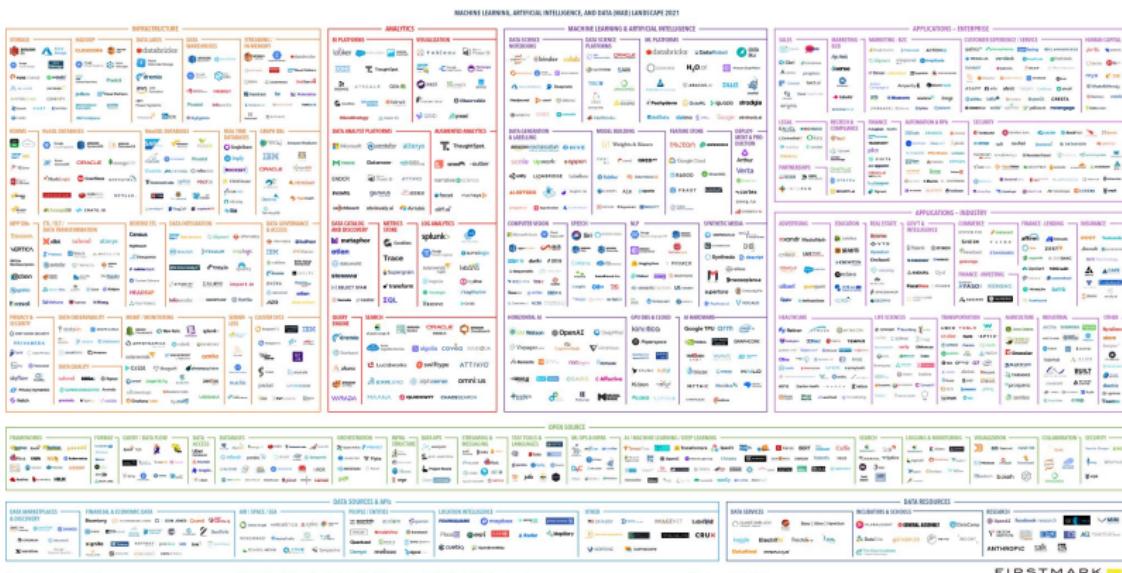
Francesca Bugiotti

CentraleSupélec

October 3, 2023

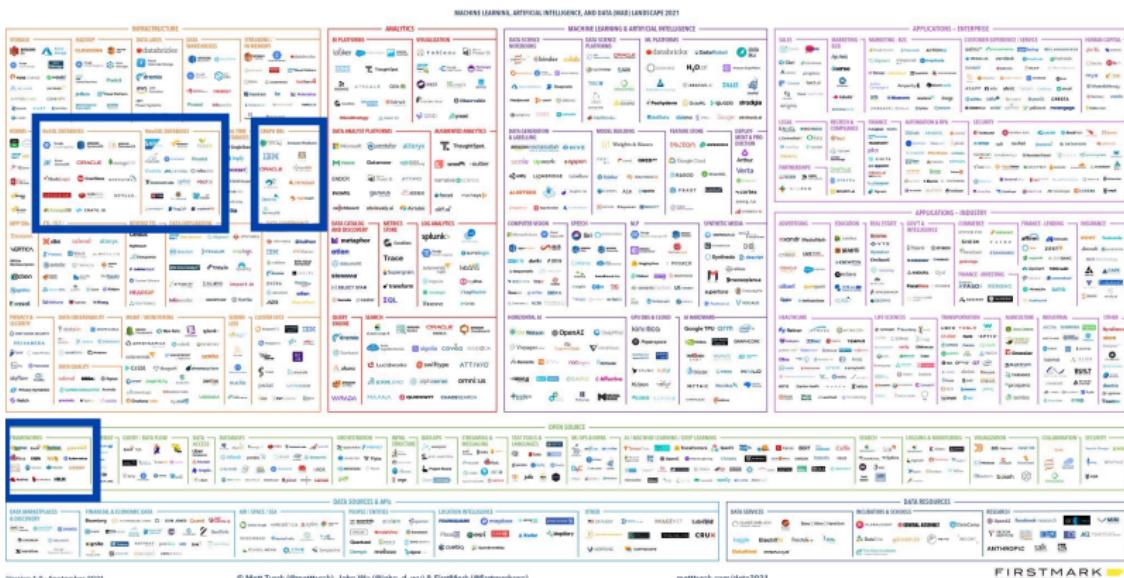


Technologies





Technologies





Objectives

- Define a distributed file system
- Recognize MapReduce programming model



Plan

1 Distributed File System



File System

When computers first came out, the information and programs were stored in punch cards



These punch cards were stored in file cabinets





File System

A **File System** defines structure and logic rules used to manage the groups of information and their names in a storage device

Long-term information storage

- Providing access for processes
- Storing large amount of information



File System

You have your 2TB disk:





File System

If the space is not sufficient?

- A bigger disk



- An external hard drive



Issues: copy data, transfer data, maintain data, etc.



Many devices implies many disks

- From two PC



- To many heterogeneous devices



Issues: copy data, transfer data, maintain data, etc.



Parallel computer vs Commodity Cluster

- **Parallel computers**
 - Single computing nodes with specialized capabilities connected by a network
- **Commodity Cluster**
 - Affordable parallel computers with an average number of computing nodes

Many jobs that share nothing can work on different data sets or parts of data sets

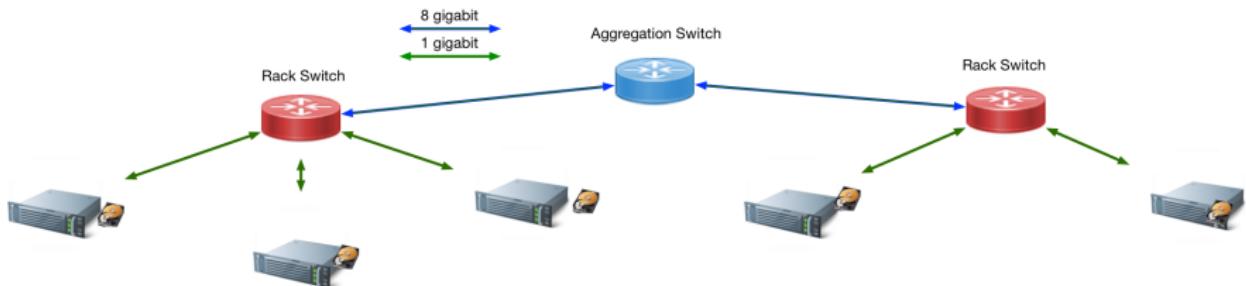
The computing in one or many nodes of those clusters is called distributed computing



Many disks

What about a system that can handle the data access and do the data distribution in a transparent way?

- Distribution across the local or wide area network



Compute nodes are **connected** by a **network**

- Typically gigabit Ethernet



Local Cluster



A cluster consists of a set of compute nodes working together:

- In a cluster you can find many compute nodes stored locally



Many disks

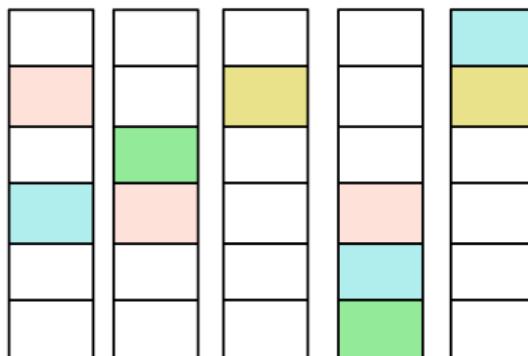
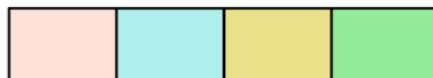
Distributed File System

- Data **distributed** across the cluster
- Data possibly **replicated** for fault tolerance



Data Distribution

Data

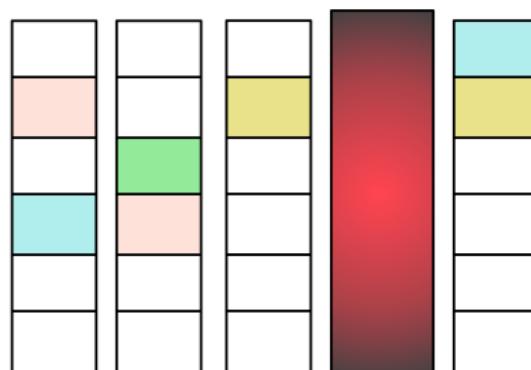
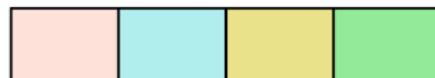


Distributed File System



Data Distribution

Data



Distributed File System



Fault tolerance

Fault tolerance: the ability to restart the computation given a failure

- Redundant data storage
- Restart of failed individual parallel jobs



Restart failed jobs

Suppose that we have a set of jobs that collaborate in order to solve a problem and belong to a main process

What can be done if one (or more) fails?

- Start all the process from the beginning
- Record checkpoint information so that you can restart just some jobs if they fail
- Make the jobs independent and restart just the failed jobs



Big Data

Big Data processing includes the features

- Multiple data types
- Large datasets



Big Data

Large volumes of data: **splitting data**

- Partitioning and placement of data in and out of computer memory
- Model to synchronize the datasets
- Replications and recovery of files

Large volumes of data: **access data**

- Fast data access
- Move the computation to data
- Scheduling of many parallel tasks
- Support data variety



Programming model

A **programming model** for Big Data should handle all these features

- Splitting large volumes of data
- Abstraction over a distributed file system
 - Handle replications
 - Synchronize datasets
- Schedule parallel tasks
- Recover from failures
- Handle different data types
- Move the computation near the data



Starting from Pasta

You come back from work on a Friday evening

You realize that you have forgotten that you have organized a party at your place that evening

What can be done?

Italian solution: prepare a pasta party!



Pasta Party



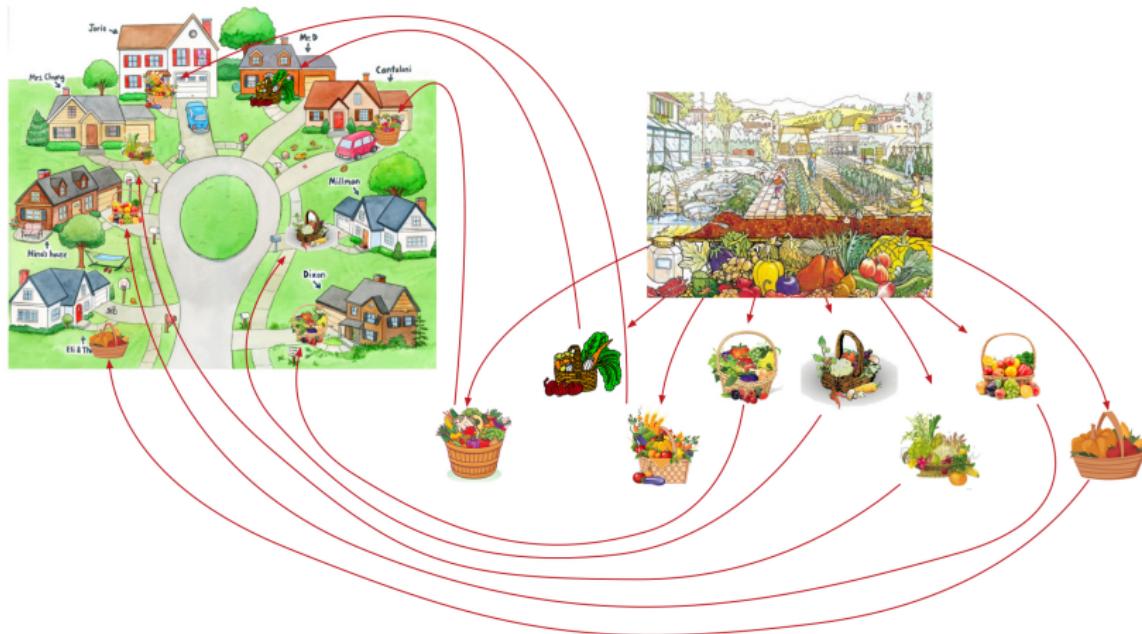


Pasta Party





Pasta Party



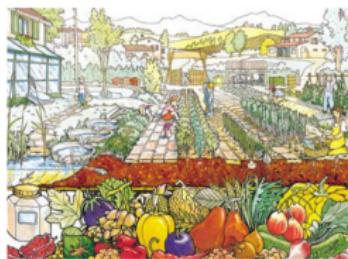


Pasta Party





Pasta Party





Pasta Party





Pasta Party





Pasta Party





Pasta Party





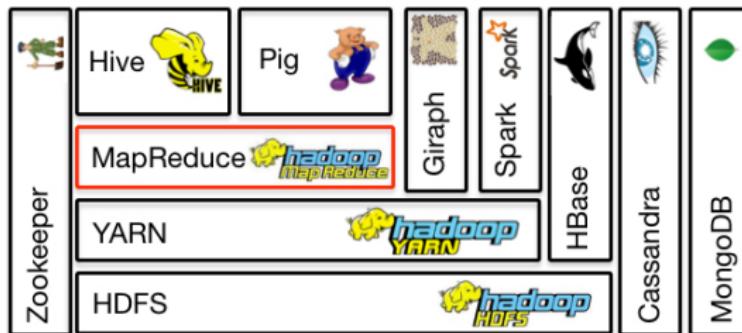
A Programming model: MapReduce

- Splitting large volumes of data
- Abstraction over a distributed file system
 - Handle replications
 - Synchronize datasets
- Schedule parallel tasks
- Recover from failures
- Handle different data types
- Move the computation near the data

Implemented by a large number of frameworks including Hadoop



MapReduce



Hadoop MapReduce [1]

With MapReduce you must create: map tasks, reduce tasks and run them



MapReduce

Traditional parallel programming

- Locks
- Semaphores
- Monitors

An incorrect usage:

- Incorrect programs
- Performance impact



Impact





MapReduce

Map

takes as input an object with a key k and a value v (k,v) and returns a bunch of key-value pairs:

$$(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)$$



MapReduce

Map

takes as input an object with a key k and a value v (k, v) and returns a bunch of key-value pairs:

$$(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)$$

- if the input set of values v_i has no keys?



MapReduce

Map

takes as input an object with a key k and a value v (k, v) and returns a bunch of key-value pairs:

$$(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)$$

- We can continue to abstract:
 - the value is also the key
 - dummy keys can be useful



MapReduce

Map

takes as input an object with a key k and a value v (k, v) and returns a bunch of key-value pairs:

$$(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)$$

Shuffle

The framework collects all the pairs with the same key k and associates with k all the values for k :

$$(k, [v_1, \dots, v_n])$$

Reduce

takes as input a key and a list of values $(k, [v_1, \dots, v_n])$ and combine them somehow



MapReduce programming

Programmers specify two functions:

- map $(k_1, v_1) \rightarrow [(k_2, v_2)]$
- reduce $(k_1, [v_1]) \rightarrow [(k_2, v_2)]$

where (k, v) denotes a (key, value) pair and $[...]$ denotes a list



MapReduce program

A MapReduce program, referred to as a job, consists of:

- Code for Map and a code for Reduce packaged together
- Configuration parameters (where the input lies, where the output should be stored)
- The input, stored on the underlying distributed file system
- Each MapReduce job is divided by the system into smaller units called tasks
 - Map tasks
 - Reduce tasks
- Input and output of MapReduce jobs are stored on the underlying distributed file system



MapReduce execution process I

- ① Some Map tasks are given each one or more chunks of data
- ② Each Map task turns the chunk into a sequence of key-value pairs
 - The way key-value pairs are produced is determined by the code written by the user for the Map function
- ③ The key-value pairs from each Map task are collected by a master controller and sorted and grouped by key (Shuffle and Sort)
- ④ The keys are divided among all the Reduce tasks, so all key-value pairs with the same key wind up at the same Reduce task
- ⑤ The Reduce tasks work **on one key at a time**, and combine all the values associated with that key in some way
 - The **way values are combined** is determined by the code written by the user for the **Reduce** function



MapReduce execution process II

- ⑥ Output: key-value pairs from each reducer are written persistently back onto the distributed file system
- ⑦ The output ends up in r files, where r is the number of reducers
 - The r files often serve as input to yet another MapReduce job



Procedure



Goflettes à la noix de coco

Préparation de l'assiette

1 C. à soupe de sucre en poudre
1 C. à soupe de farine de riz
1 C. à soupe de beurre fondu
1 C. à soupe de lait
1 C. à soupe de noix de coco râpée

1. Préchauffer le four à 180°C.
2. Dans un bol, mélanger le sucre et la farine de riz. Ajouter les noix de coco râpées et bien mélanger les deux ingrédients.
3. Incorporer progressivement la farine et le beurre fondu jusqu'à ce qu'il n'y ait plus trace de farine.
4. Préchauffer une poêle à fond plat avec un peu d'huile végétale.
5. Ajouter la pâte dans la poêle et cuire jusqu'à ce que la pâte soit dorée et croustillante.

Recette pour 10 personnes

Préparation de l'assiette

1. Couper la pomme en dés et la faire cuire dans une poêle avec de la crème.
2. Faire cuire les goflettes.
3. Garnir l'assiette avec les goflettes et la crème pomme.

Faisselle (250g)

La dégustation

Fromage à la crème

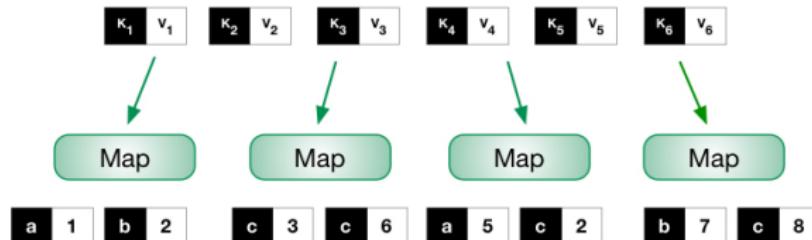
Crème fraîche

Confiture

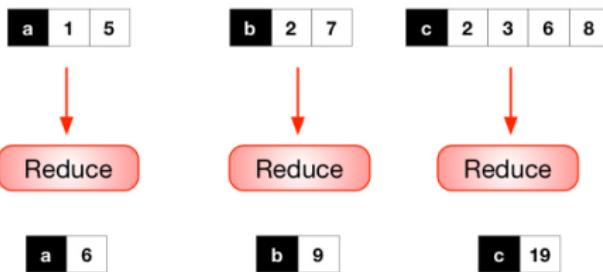




MapReduce



Shuffle and sort: aggregate values by keys





Counting words I

Problem:

counting the number of occurrences of each word in a collection of documents

- **Input**: a repository of documents, each document is an element
- **Map**: reads a document and emits a sequence of key-value pairs where keys are words of the documents and values are equal to 1:

$$(w_1, 1), (w_2, 1), \dots, (w_n, 1)$$

- **Shuffle**: groups by key and generates pairs of the form

$$(w_1, [1, 1, \dots, 1]), \dots, (w_n, [1, 1, \dots, 1])$$

- **Reduce**: adds up all the values and emits:



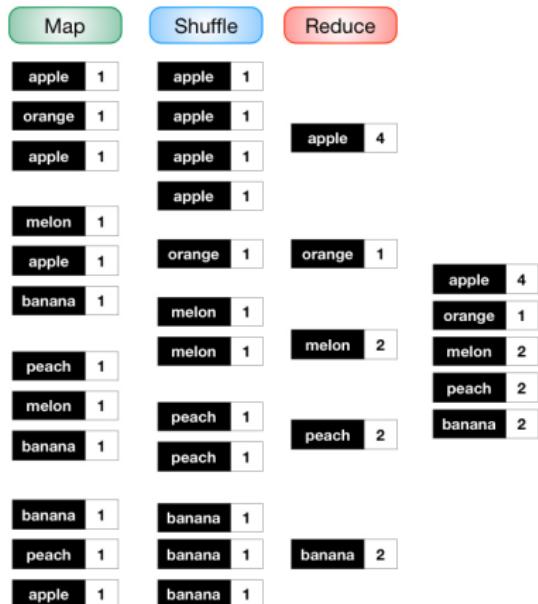
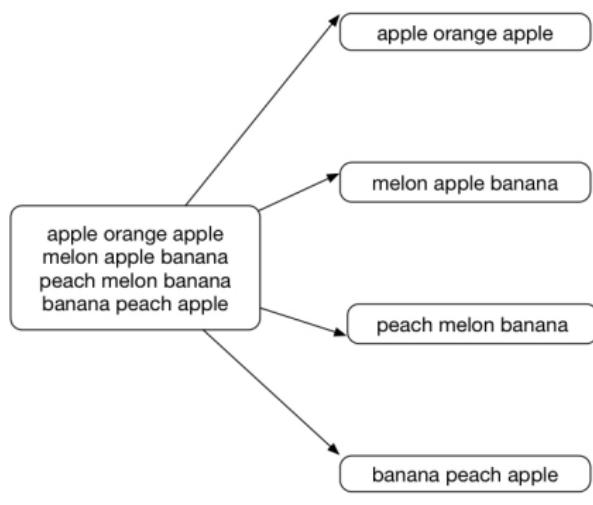
Counting words II

$$(w_1, k), \dots, (w_n, l)$$

- **Output:** (w, m) pairs, where w is a word that appears at least once among all the input documents and m is the total number of occurrences of w among all those documents



MapReduce





Implementation - pseudocode

```
Map( String docID, String text):
    for each word w in text:
        Emit(w, 1);

Reduce( String term, counts []):
    int sum = 0;
    for each c in counts:
        sum += c;
    Emit(term, sum);
```



Implementation Python

```
for line in sys.stdin:  
    line = line.strip()  
    words = line.split()  
    for word in words:  
        print '%s\t%s' % (word, 1)
```



Exercise 1

You have as input a big file that is stored in a distributed environment. This big files contains a list of text messages that have been sent in a company context. For making this exercise easy consider that:

- each line of the file contains one and only one message.
- there are not additional information provided (sender, receiver, time, etc.)
- privacy is not important then no cryptography is applied

Provide a map-reduce algorithm that provides one of the messages of maximum length.

- A kaggle dataset that can be used for trying your code



Combiners

When we can associate with the Reduce a function that is **associative** and **commutative**, we can push some of what the reducers do to the Map tasks

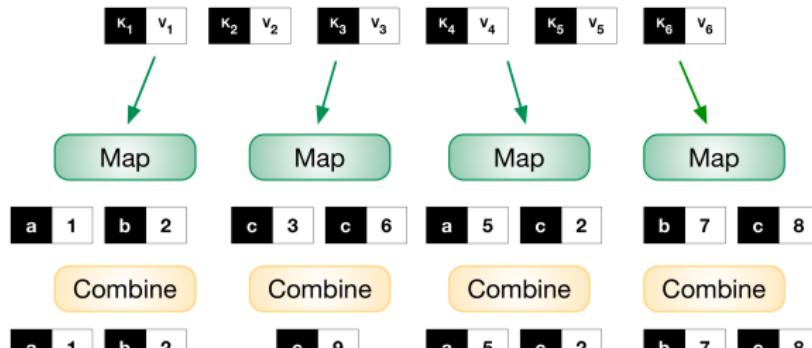
- In this case we also apply a combiner to the Map function
- In many cases the same function can be used for combining as the final reduction
- Shuffle and sort are still necessary!

Advantages:

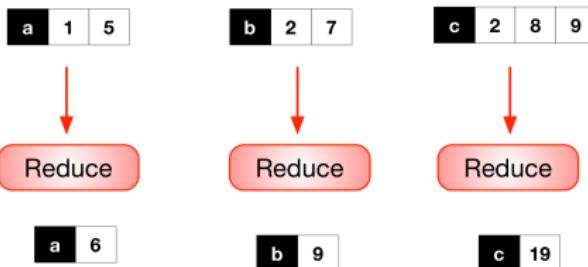
- It reduces the amount of **intermediate** data
- It reduces the **network traffic**



MapReduce

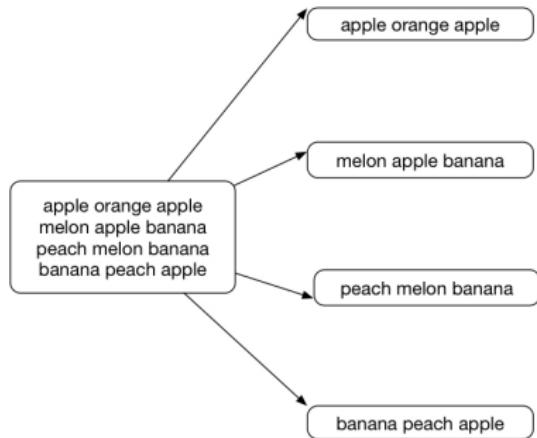


Shuffle and sort: aggregate values by keys





MapReduce



Map	Combine	Shuffle	Reduce
apple 1	apple 2	apple 2	
orange 1	orange 1	apple 1	
apple 1		apple 1	apple 4
melon 1	melon 1	orange 1	orange 1
apple 1	apple 1	melon 1	
banana 1	banana 1	melon 1	
		peach 1	
peach 1	peach 1	peach 1	
melon 1	melon 1	peach 1	
banana 1	banana 1	peach 1	peach 2
banana 1	banana 1	banana 1	
peach 1	peach 1	banana 1	banana 2
apple 1	apple 1	banana 1	



Counting words with combiners I

Problem:

counting the number of occurrences for each word in a collection of documents

- **Input:** a repository of documents, each document is an element
- **Map:** reads a document and emits a sequence of key-value pairs where keys are words of the documents and values are equal to 1:

$$(w_1, 1), \dots, (w_n, 1)$$

- **Combiner:** groups by key, adds up all the values and emits:

$$(w_1, i), \dots, (w_n, j)$$

- **Shuffle:** groups by key and generates pairs of the form



Counting words with combiners II

$(w_1, [1, 1, \dots, 1]), \dots, (w_n, [1, 1, \dots, 1])$

- **Reduce:** adds up all the values and emits:

$(w_1, k), \dots, (w_n, l)$

- **Output:** (w, m) pairs, where w is a word that appears at least once among all the input documents and m is the total number of occurrences of w among all those documents



Exercise 2

You have as input a big file that is stored in a distributed environment. This big files contains a list of text messages that have been sent in a company context.

For making this exercise easy consider that:

- each line of the file contains one and only one message.
- there are not additional information provided (sender, receiver, time, etc.)
- privacy is not important then no cryptography is applied

Provide a map-reduce algorithm that provides **one of the messages of maximum length**.

- A kaggle dataset that can be used for trying your code



Exercise 3

You have as input a big file that is stored in a distributed environment. This big files contains a list of text messages that have been sent in a company context. For making this exercise easy consider that:

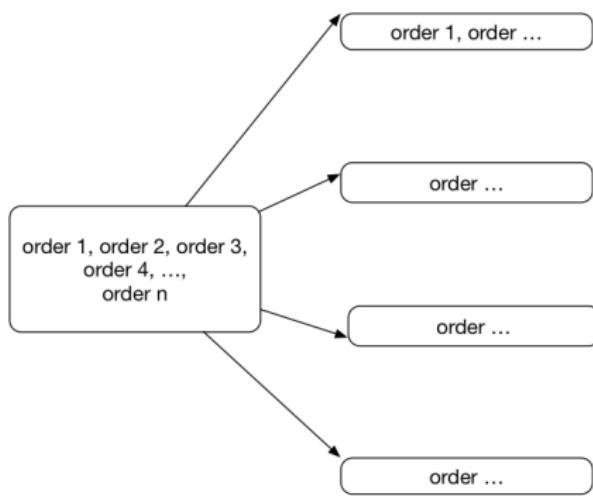
- each line of the file contains one and only one message.
- there are not additional information provided (sender, receiver, time, etc.)
- privacy is not important then no cryptography is applied

Provide a map-reduce algorithm that provides the **average length of the exchanged messages**.

- A kaggle dataset that can be used for trying your code



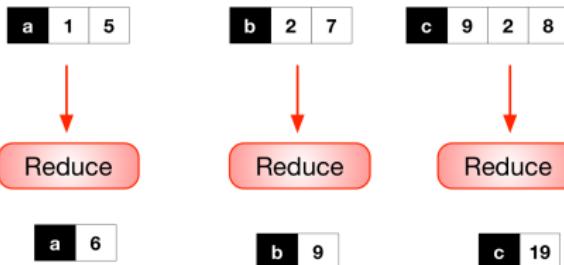
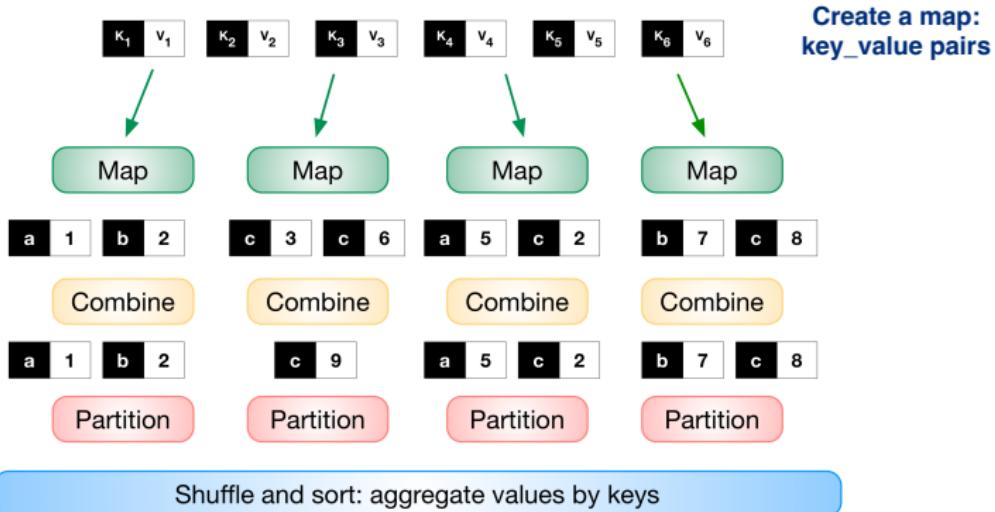
Different combiners and reducers



Map	Shuffle	Reduce
disk u1	disk u1	
mouse u3	disk u1	disk 3
disk u5	disk u5	
	disk u7	
	usb u3	usb 1
	screen u4	screen 2
	laptop u7	laptop 1
	usb u3	cable 1
	screen u2	laptop 1
	cable u4	printer 1
	printer u9	mouse 2
	mouse u1	
		mouse 2



Different combiners and reducers





MapReduce

Partitioner

- **divides** up the intermediate key space
 - **assigns** intermediate key-value pairs to reducers
 - n partitions and n reducers
-
- Assigns approximately the same number of keys to each reducer
 - Considers the key and ignores the value
 - Word frequency in inverted indexes



More operations

Map and reduce are transformations that work on a single list of keys. We can play with collections sharing the same keys:

- **cross product**, or a cartesian product
- **match** or join
- **co-group**

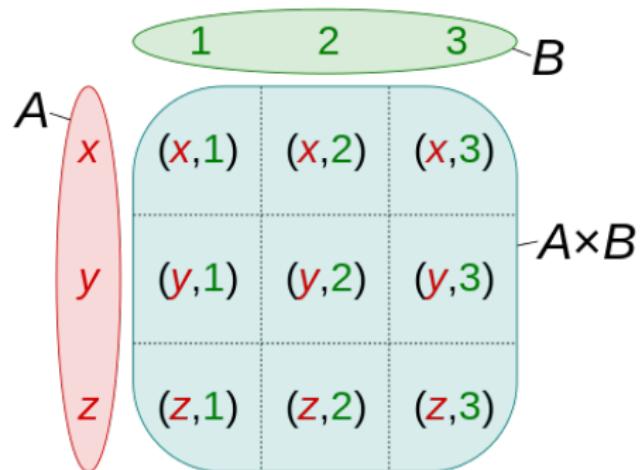


Cartesian Product

The Cartesian product of two sets A and B, denoted as $A \times B$, is an operation that produces the set of all the ordered pairs (a, b) where a is in A and b is in B.

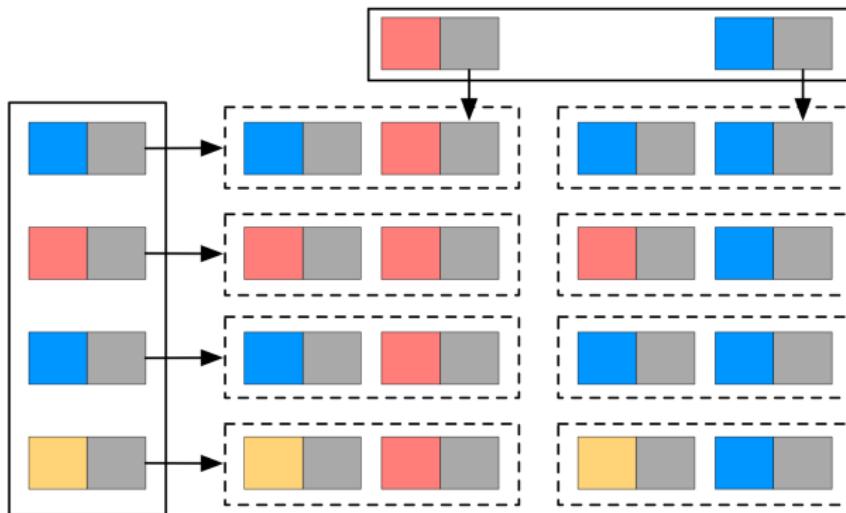


Cartesian Product





Cartesian Product



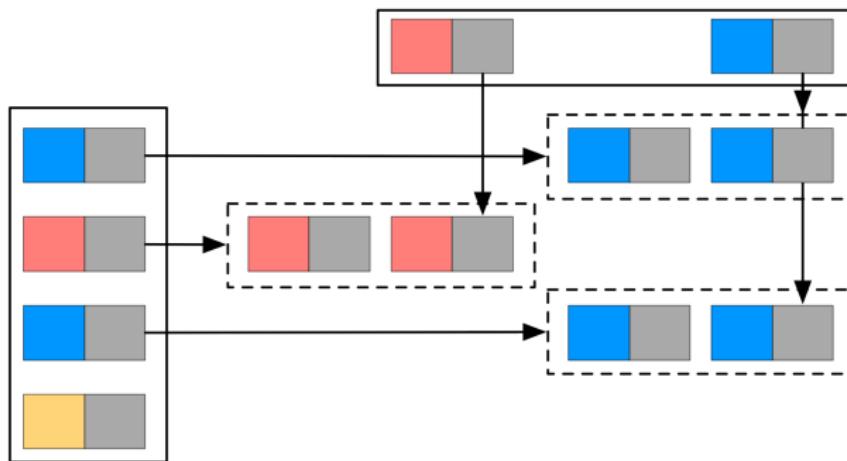


Match

The match operation compares values for identifying equal instances in two sets.



Match



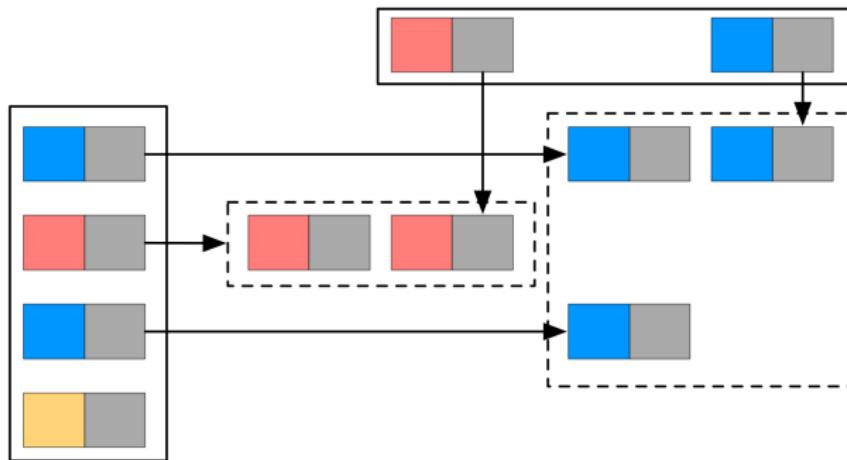


Co-group

The co-group operations puts together equal elements.



Co-group





Summary

MapReduce hides complexities of parallel programming and greatly simplifies building parallel applications

- Data gets reduced to a smaller set at each step

Not good if:

- Data is frequently changing
- Map and reduce are dependent
- You need intermediate results



References |



Hadoop Map Reduce.

<http://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>.

Accessed 2016.