



Big Data, Techniques and Platforms

Class 2/8: Data Distribution

Francesca Bugiotti

CentraleSupélec

October 3, 2023



Objectives

- Data distribution models
- Consistency in distributed databases
- Lesson 5: Availability in distributed databases
- Lesson 5: CAP theorem



Plan

1 Summary

2 Scalability

3 Data Distribution



Summary

Data Science

- **Data**
 - Common devices and specialized devices
 - New data sources producing torrent of data
- **Cloud Computing**
 - Computing anywhere and anytime
 - On-demand Computing



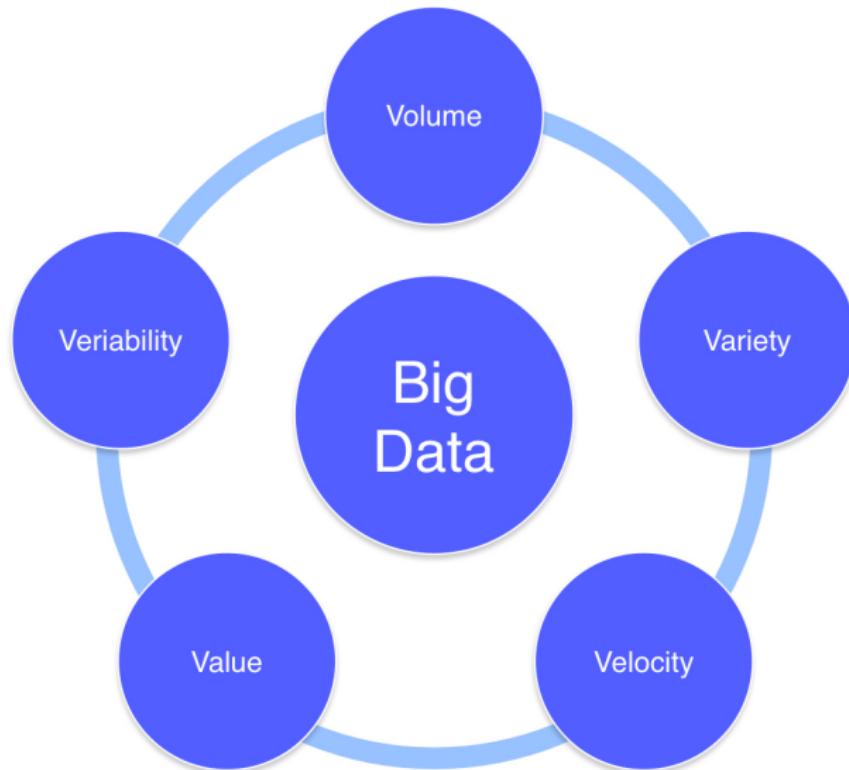
Big data is not just “Big”

Combining data from different sources

- **Machines**
 - real-time data, sensor data, trackers, streaming, etc.
- **People**
 - social media, personal documents, etc.
- **Organizations**
 - structured knowledge bases, data-warehouses, etc.



Big Data: the 5 Vs





MODERN DATA SCIENTIST

Data Scientist, the sexiest job of the 21th century, requires a mixture of multidisciplinary skills ranging from an intersection of mathematics, statistics, computer science, communication and business. Finding a data scientist is hard. Finding people who understand who a data scientist is, is equally hard. So here is a little cheat sheet on who the modern data scientist really is.

MATH & STATISTICS

- ★ Machine learning
- ★ Statistical modeling
- ★ Experiment design
- ★ Bayesian inference
- ★ Supervised learning: decision trees, random forests, logistic regression
- ★ Unsupervised learning: clustering, dimensionality reduction
- ★ Optimization: gradient descent and variants

DOMAIN KNOWLEDGE & SOFT SKILLS

- ★ Passionate about the business
- ★ Curious about data
- ★ Influence without authority
- ★ Hacker mindset
- ★ Problem solver
- ★ Strategic, proactive, creative, innovative and collaborative



PROGRAMMING & DATABASE

- ★ Computer science fundamentals
- ★ Scripting language e.g. Python
- ★ Statistical computing packages, e.g., R
- ★ Databases: SQL and NoSQL
- ★ Relational algebra
- ★ Parallel databases and parallel query processing
- ★ MapReduce concepts
- ★ Hadoop and Hive/Pig
- ★ Custom reducers
- ★ Experience with xaaS like AWS

COMMUNICATION & VISUALIZATION

- ★ Able to engage with senior management
- ★ Story telling skills
- ★ Translate data-driven insights into decisions and actions
- ★ Visual art design
- ★ R packages like ggplot or lattice
- ★ Knowledge of any of visualization tools e.g. Flare, D3.js, Tableau



Bakery





Baker - Chef





Baker - Chef





Ingredients





Stockage Utils





Utils





Recipes



Gâteaux à la noix de coco

la pâte :
150 g de beurre
150 g de sucre
1 verre de sucre
1 œuf
1 cuillère à soupe d'eau
1 cuillère à soupe de levure chimique
1 cuillère à soupe de sucre en poudre
1 cuillère à soupe de noix de coco râpée

la garniture :
1 paquet de sucre à sucer
1 C. à soupe de sucre en poudre
1 cuillère à soupe de noix de coco râpée
Confiture

Préparation :
1. Dans un bol, faire fondre le beurre et le sucre au bain-marie. Ajouter l'œuf et bien mélanger le tout.
2. Ajouter la levure et la noix de coco râpée et bien mélanger.
3. Ajouter la pâte à la garniture et bien mélanger.
4. Mettre la pâte dans un moule et cuire au four à 180°C pendant 20 minutes.

Recette de la garniture :

1. Faire fondre le beurre et le sucre à feu doux puis ajouter la noix de coco râpée.
2. Ajouter la confiture et bien mélanger.
3. Verser la garniture dans une forme à gâteau et enfourner au four à 180°C pendant 10 minutes.

Étapes de préparation :

1. Préparation de la pâte
2. Préparation de la garniture
3. Cuissage du gâteau
4. Garniture
5. Finalisation

gâteau incroyable
POMME CHOCO-CARAMEL

Préparation :
1. Préchauffer le four à 180°C.
2. Dans un bol, mélanger la farine, la levure et la poudre d'amande.
3. Ajouter le beurre et les noix de coco râpées, puis mélanger jusqu'à obtention d'une pâte homogène.
4. Former une boule et la déposer sur une plaque recouverte de papier sulfurisé. Laisser reposer 15 minutes.

Recette de la garniture :

1. Faire fondre le beurre et ajouter la crème liquide et le sucre à sucer.
2. Ajouter la crème de cacao et bien mélanger.
3. Ajouter les noix de coco râpées et bien mélanger.

Préparation de la garniture :

1. Préparation de la pâte
2. Cuissage de la pâte
3. Préparation de la garniture
4. Garniture





Procedure



pâtisserie
POMME CHOCO-CARAMEL





Procedure



Goflettes à la noix de coco

pour 250 g de beurre
- 1 verre de sucre glace + 1 cuillère à soupe de sucre de canne
- 1 pincée de sel
- 1 cuillère à soupe de levure chimique (C & C ou celle de biscuits) + 1 verre de lait tiède

avant un morceau de beurre court et mélanger le sucre, le sel et la levure dans un bol. ajouter les œufs et mélanger jusqu'à ce que le pâte commence à prendre forme. faire une pâte moelleuse.

2. Préparer des goflettes en utilisant une poche à pâtisserie garnie d'une poche à sucre.

3. Faire cuire les goflettes dans un four préchauffé à 180°C pendant 10 minutes.

4. Bien réserver sur une grille pour qu'ils refroidissent complètement.

Felicie (250g)
La déclinaison
Fruit jellée et crème vanille
Goflettes

Recette

Préparation

Présentation

Conseils





Procedure





In our context

- Data Science
- Data
- Data Scientist
- Algorithms: K-means, PageRank, etc.
- Databases
- Map-Reduce
- Hadoop, Spark
- Produce Value



Plan

1 Summary

2 Scalability

3 Data Distribution



Scalability

Fix the terminology and the scenarios:

- Scale up
- Scale out



Single Server





Single Server

How is data organized?

No distribution

Data is on a **single machine**

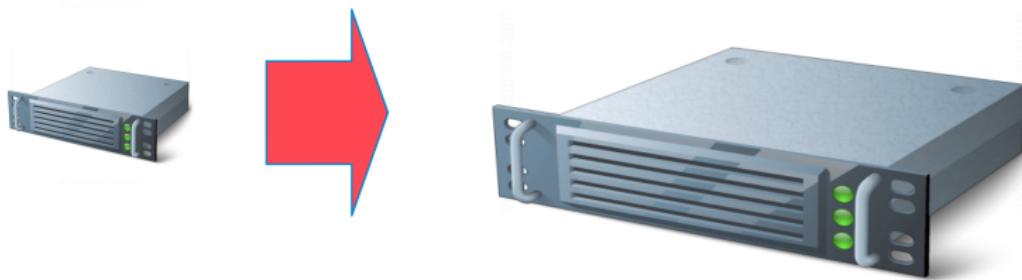


Scale up



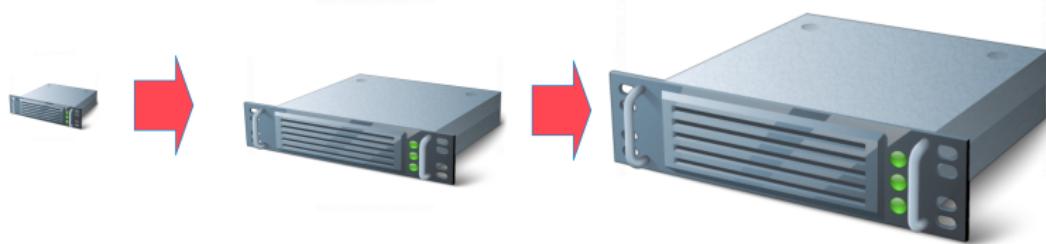


Scale up





Scale up





Single Server - Scalability

How do we achieve scalability?

Scenario - Single Server

Methodology - **Scale up:**

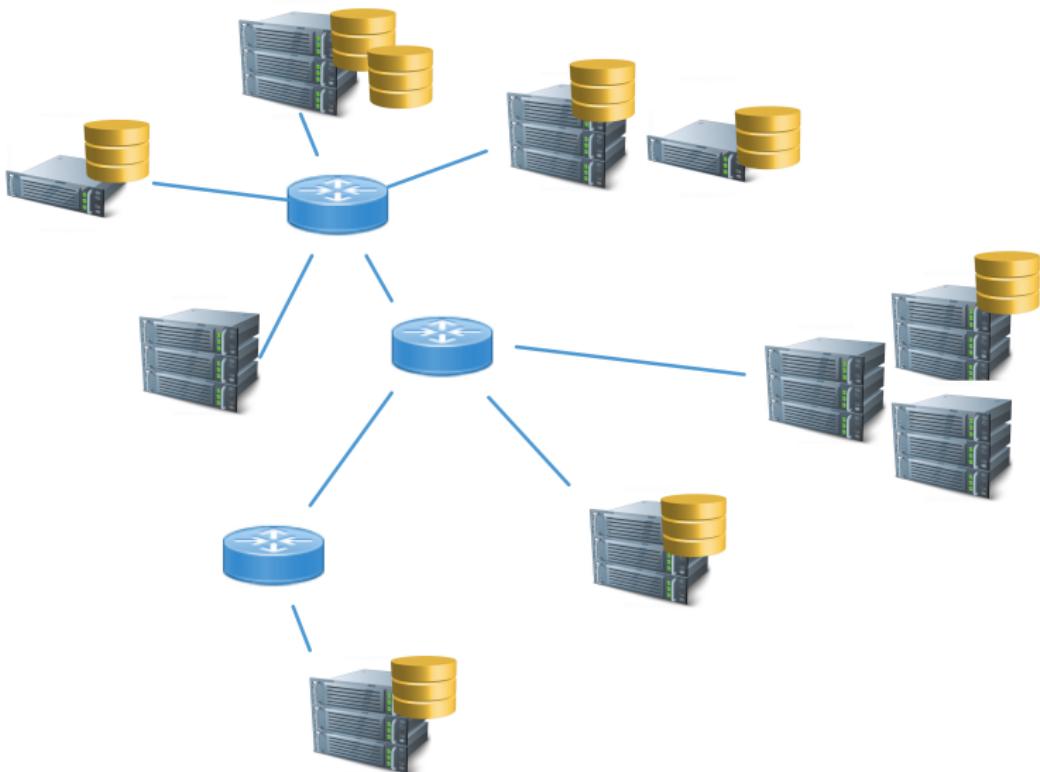
- Add new computation capability **if needed** and/or **if possible**

The single machine handles all the reads and writes to data

- Hint for the future: fits well with graph databases



Multiple servers





Multiple servers

How is data organized?

Data is distributed

Data is on **multiple nodes** of a cluster

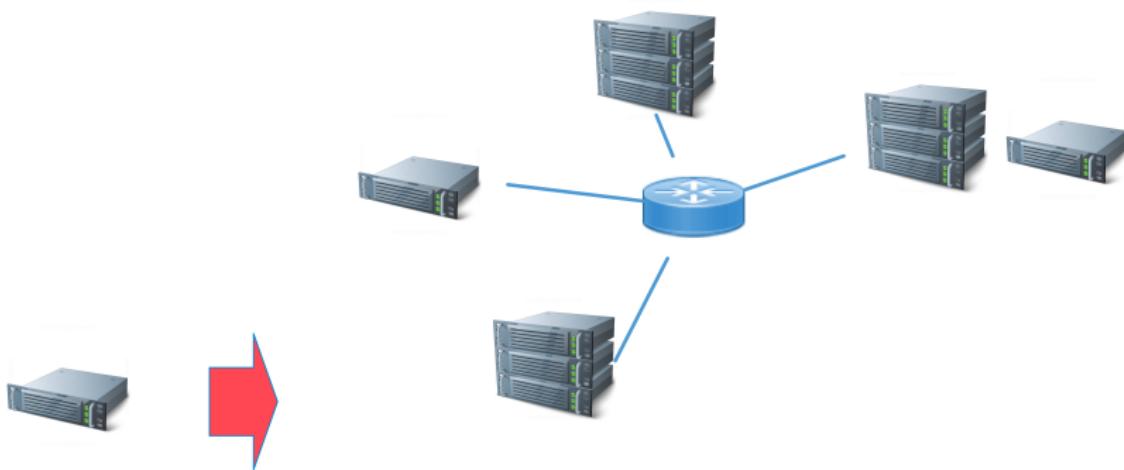


Scale Out



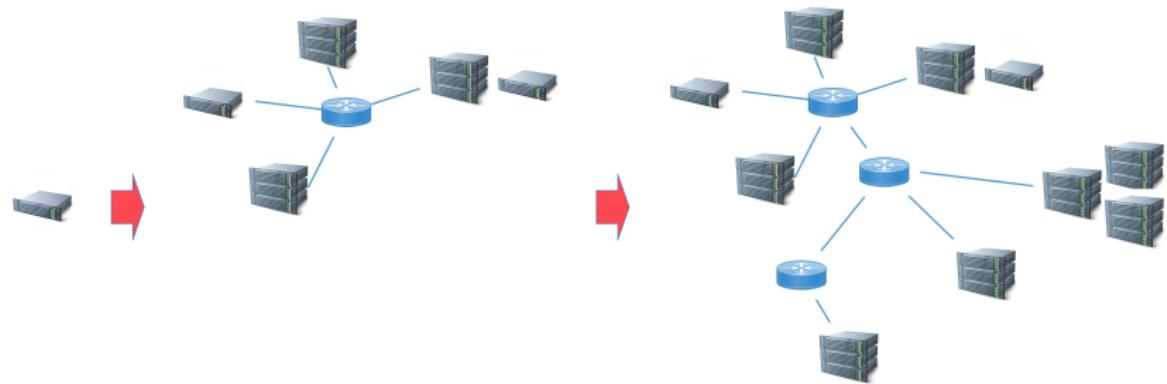


Scale Out





Scale Out





Multiple Servers - Scalability

How do we achieve scalability?

Scenario - Multiple Servers

Methodology - **Scale out:**

- Run the database on a cluster of servers
- Add new nodes when needed

Related issues

- ① How to distribute (shard) data across the nodes
- ② How and if replicate data
- ③ How to assign read/write operations
 - Focus on read/write operations



Plan

1 Summary

2 Scalability

3 Data Distribution



Data Distribution

- Data distribution models:
 - Single server
 - Multiple servers
- Orthogonal aspects of data distribution:
 - Sharding: different data on different nodes
 - Replication: the same data copied over multiple nodes



Sharding

Different parts of data into **different nodes** and each node does its own reads and writes

- **Horizontal scalability**
- Ideal case: different users all talking to different server nodes
- Data accessed together on the same node
 - aggregate data!
 - How?
 - This is a challenge!!!
- **Pros:** it can improve both reads and writes
- **Cons:** clusters use less reliable machines - resilience decreases

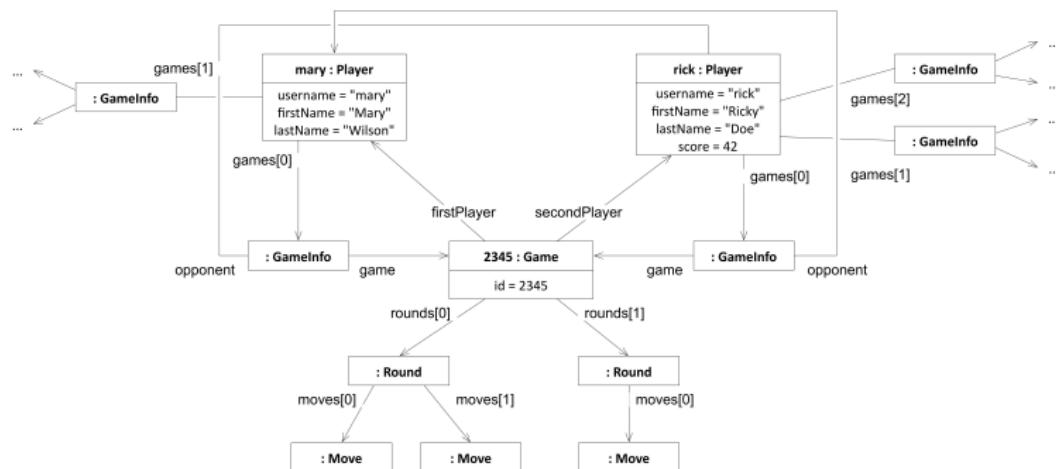


Aggregates

- Data as units that have a complex structure
 - More structure than just a set of tuples
 - A complex record with fields inside
 - A record with a composite access structure
- Aggregates in Domain Driven Design [2]:
 - A collection of related objects that we treat as a unit
 - A unit for data manipulation and management of consistency
- Advantages of aggregates:
 - easier for programmer to work with
 - easier for database systems to handle operations on a cluster

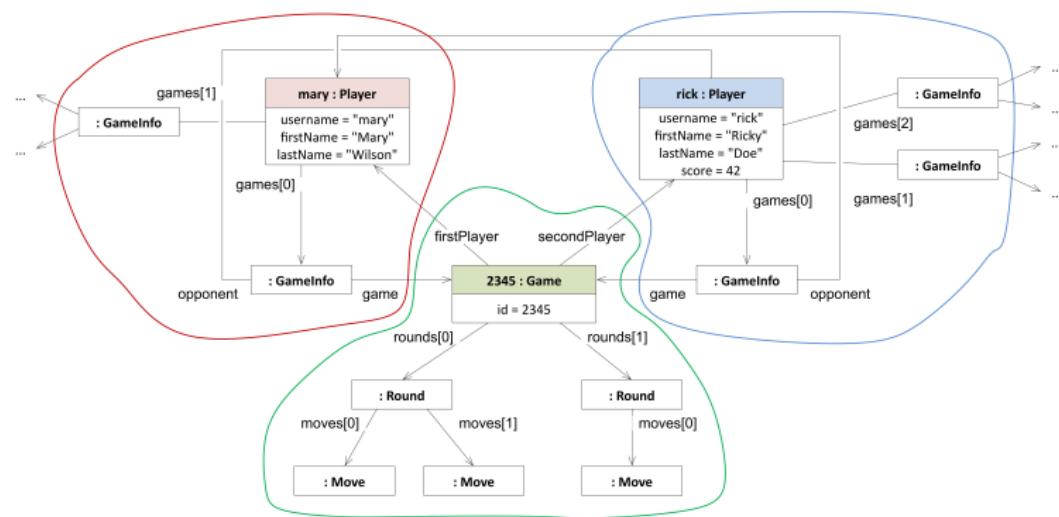


Aggregates [1]



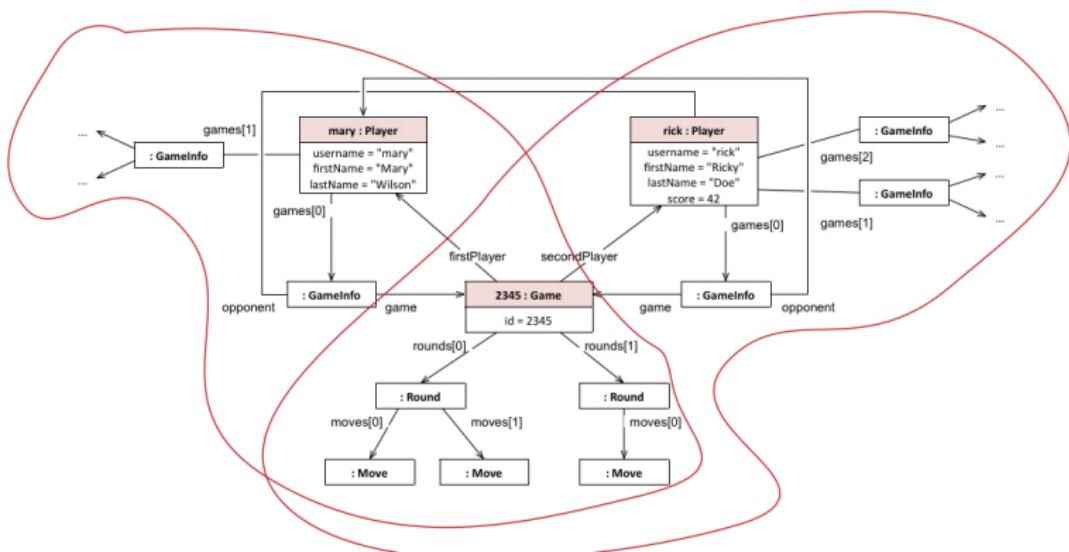


Aggregates





Aggregates





Design Strategy

- No universal strategy for designing aggregates boundaries
- The design depends on **data manipulation**
 - Access on a single round
 - Access on a player with all his games and rounds
- Context specific:
 - Some applications will run on a data representation
 - Some others on another data representation
- Focus on the **Data Access Unit**
 - very good for clusters
 - not easy to change the data interaction strategy



Sharding

Main rules of sharding:

- ① Place data close to where it's accessed
 - Games played by Californian players: data in the western US data center
- ② Try to keep the load even
 - All nodes should get equal amounts of the load
- ③ Put together aggregates that may be read in sequence
 - Rounds of the same game in the same node



Horizontal scalability

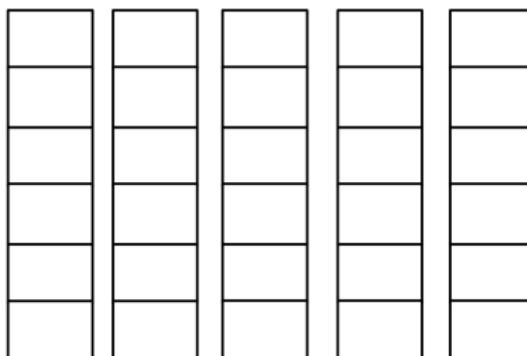
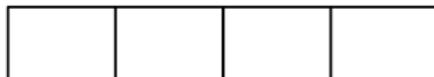
Build your database on a cluster and distribute data across the cluster

How?



Horizontal scalability

Data

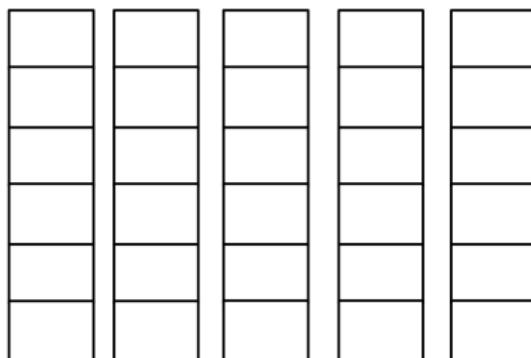


Distributed Database



Horizontal scalability

Data

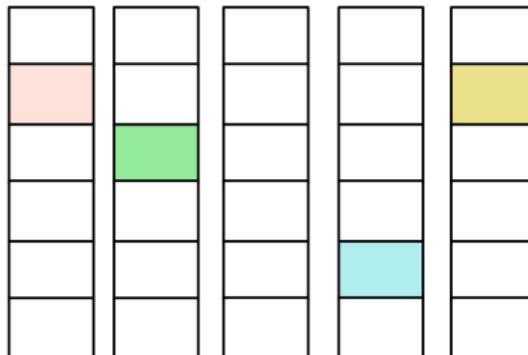


Distributed Database



Horizontal scalability

Data



Distributed Database



Horizontal scalability

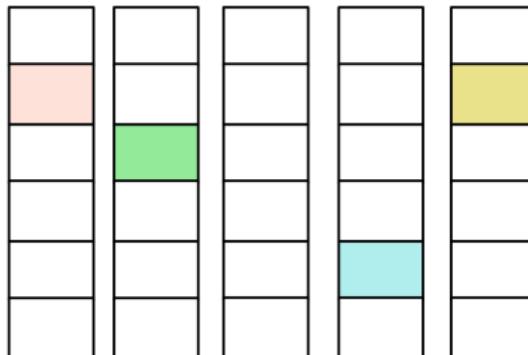
- How to divide data:
 - think about the aggregates
 - think about how to store aggregates together

Sometimes it is necessary to rethink aggregates



Horizontal scalability

Data



Distributed Database



More on Distribution

- Sharding is particularly valuable for performance
- Sharding does little for applications that have a lot of writes



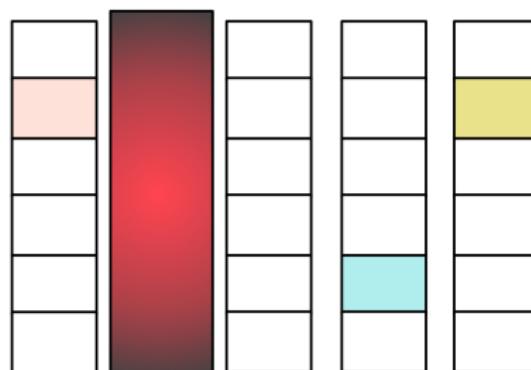
Availability

- A node failure makes pieces of data unavailable
 - Maybe not for all the users
- It is like having a single server
- Moreover clusters usually try to use less reliable machines



Availability

Data

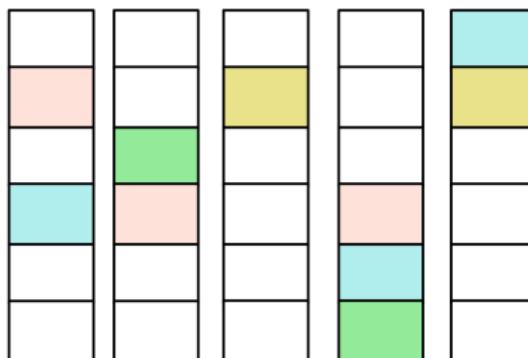


Distributed Database



Availability - Data Replication

Data

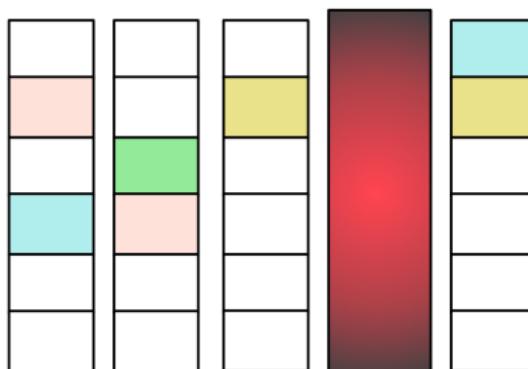


Distributed Database



Availability - Data Replication

Data



Distributed Database



Data Replication

Two forms of data replication:

- Master-slave
- Peer to peer

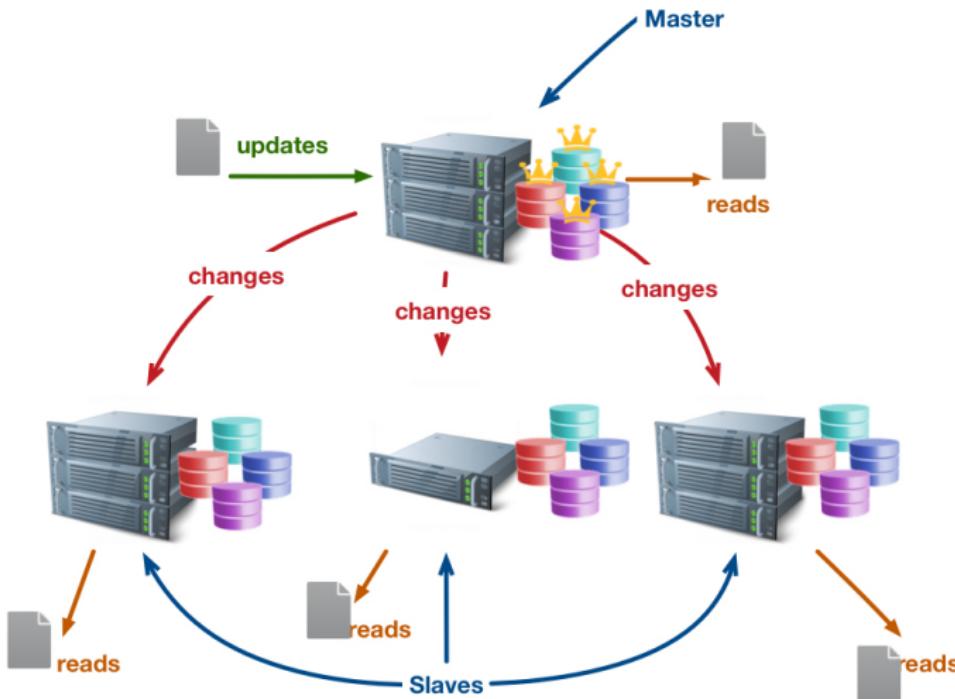


Master-Slave Replication



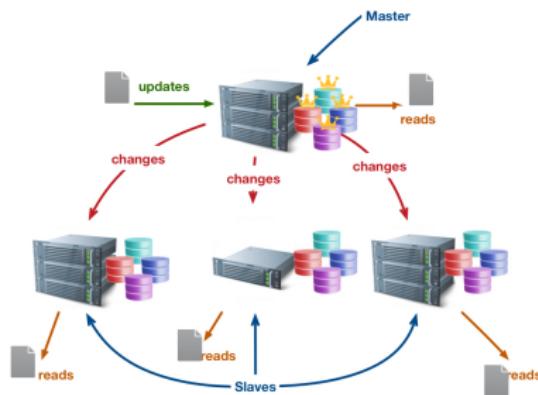


Master-Slave Replication





Master-Slave Replication

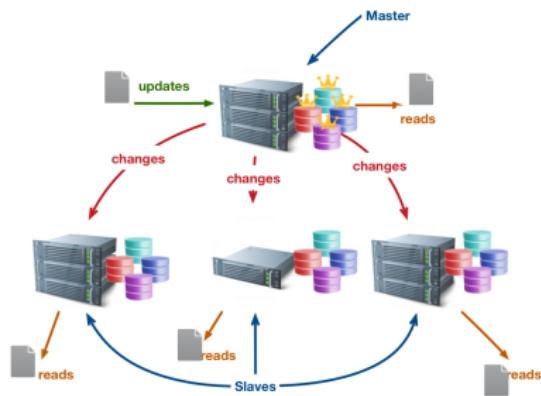


Master:

- is the authoritative source for the data
- is responsible for processing any data update
- can be appointed manually or automatically



Master-Slave Replication



Slaves:

- A replication process synchronizes the slaves with the master
- After a failure of the master, a slave can be appointed as new master very quickly



Master-slave analysis

Pros

- More read requests:
 - Add more slave nodes
 - Ensure that all read requests are routed to the slaves
- Should the master fail, the slaves can still **handle read requests**
- Good for datasets with a read-intensive dataset

Cons

- The master is a **bottleneck**
 - Limited by its ability to process updates and to pass those updates on
 - Its failure does eliminate the ability to handle writes until:
 - The master is restored or
 - A new master is appointed
- Inconsistency due to **slow propagation** of changes to the slaves
- Bad for datasets with heavy write traffic

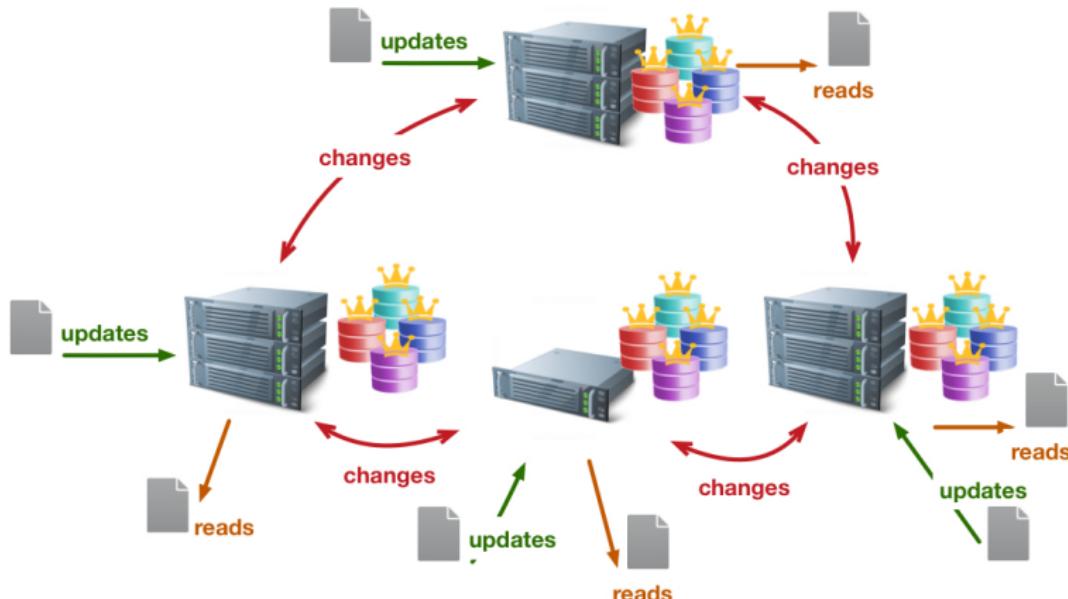


Peer-to-Peer Replication



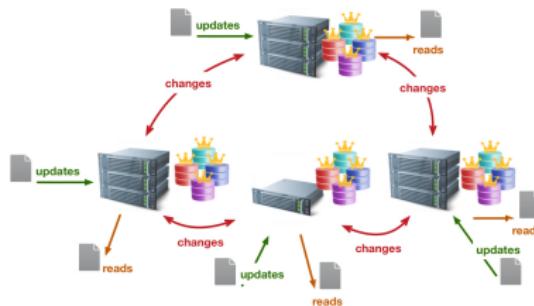


Peer-to-Peer Replication





Peer-to-Peer Replication



- All the replicas have equal weight: they can all accept writes
- The loss of any of them doesn't prevent access to the data store



Peer-to-Peer analysis

Pros

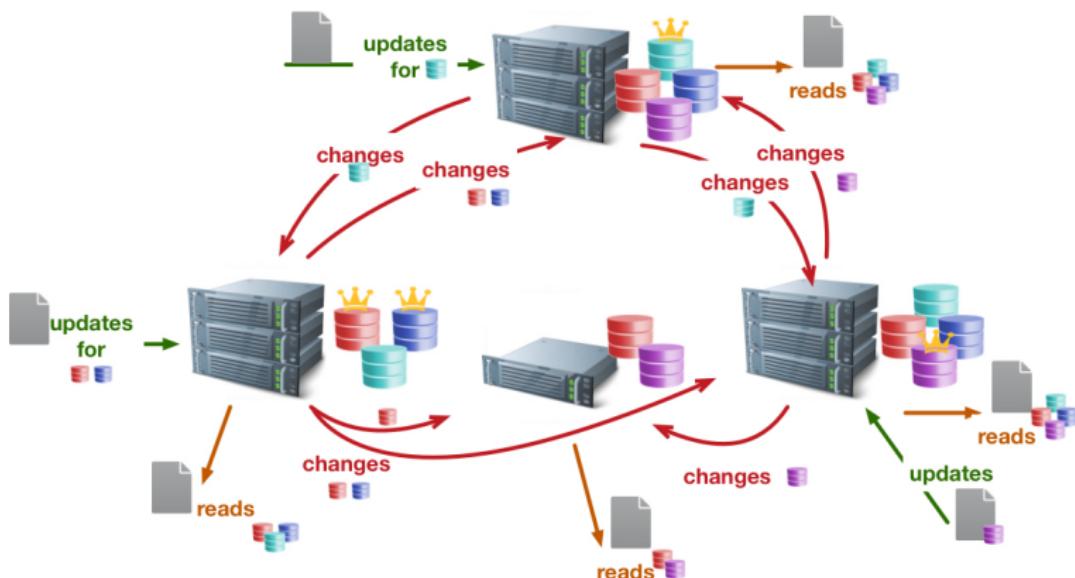
- Ride over node failures without losing access to data
- Easy to add nodes to improve performance

Cons

- Inconsistency
 - Slow propagation of changes to copies on different nodes
 - Inconsistencies on read lead to problems but are relatively transient
 - Two people can update different copies of the same record stored on different nodes at the same time:
 - a write-write conflict
 - Inconsistent writes are forever

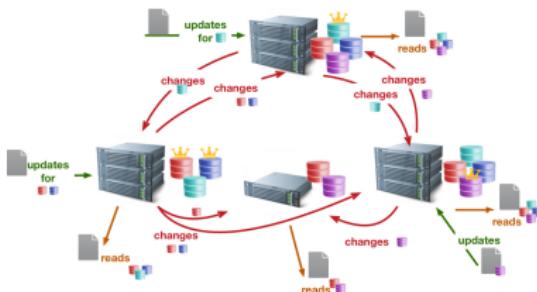


Sharding + Replication: Master-Slave





Sharding + Replication: Master-Slave

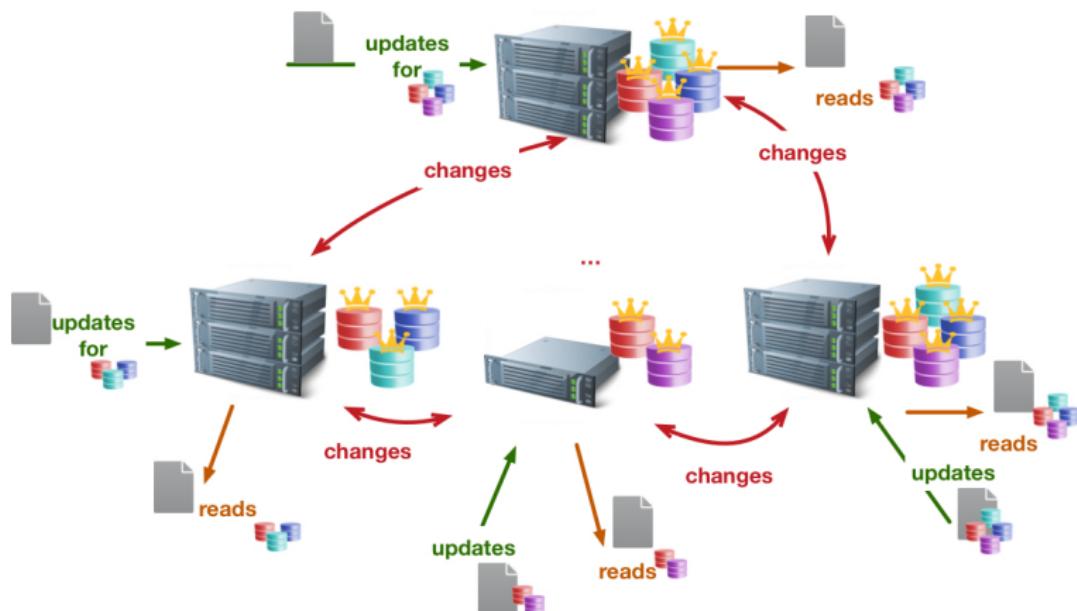


Master-Slave

- We have multiple masters, but each data only has a single master
- Two schemes:
 - A node can be a master for some data and a slave for some others
 - Nodes are dedicated for master or slave duties

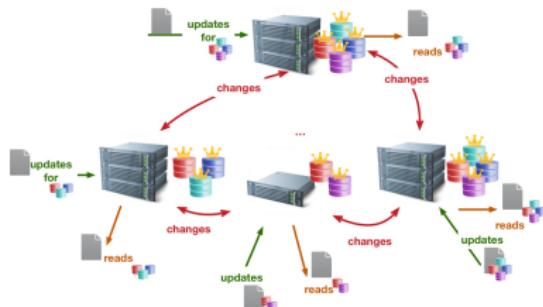


Sharding + Replication: Peer-to-Peer





Sharding + Replication: Peer-to-Peer



Peer-to-Peer

- Usually each shard is replicated on three nodes
- A common strategy for column-family databases



Summary

There are two styles of **distributing data**:

Sharding

- Distributes different data across multiple servers
 - Each server acts as the single source for a subset of data



Summary

There are two styles of **distributing data**:

Sharding

- Distributes different data across multiple servers
 - Each server acts as the single source for a subset of data

Replication

- Copies data across multiple servers
 - Each bit of data can be found in multiple places



Summary

There are two styles of **distributing data**:

Sharding

- Distributes different data across multiple servers
 - Each server acts as the single source for a subset of data

Replication

- Copies data across multiple servers
 - Each bit of data can be found in multiple places

A system may use either or both techniques



Summary

Data **replication** comes in two forms:

Master-slave replication

- makes one node the authoritative copy (master)
- the master handles writes
- slaves synchronize with the master and may handle reads



Summary

Data **replication** comes in two forms:

Master-slave replication

- makes one node the authoritative copy (master)
- the master handles writes
- slaves synchronize with the master and may handle reads

Peer-to-peer replication

- allows writes to any node
- the nodes coordinate to synchronize their copies of the data



Summary

Analysis

- Master-slave replication reduces the chance of update conflicts
- Peer-to-Peer replication avoids loading all writes onto a single point of failure



References I

-  Francesca Bugiotti, Luca Cabibbo, Paolo Atzeni, and Riccardo Torlone.
Database design for nosql systems.
In *Conceptual Modeling - 33rd International Conference, ER 2014, Atlanta, GA, USA, October 27-29, 2014. Proceedings*, pages 223–231, 2014.
-  Eric Evans.
Domain-Driven Design: Tacking Complexity In the Heart of Software.
Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2003.