



Big Data, Techniques and Platforms

Class 4/8: NoSQL

Francesca Bugiotti

CentraleSupélec

October 26, 2023



Objectives

- Understand what is a Data Model
- Review the Relational Data Model
- Understand what is a NoSQL database
- Assign its own category to a NoSQL database



Plan

- 1 Summary
- 2 Relational Data Model
- 3 NoSQL
- 4 Data Modeling in NoSQL



Ingredients





Stockage Utils



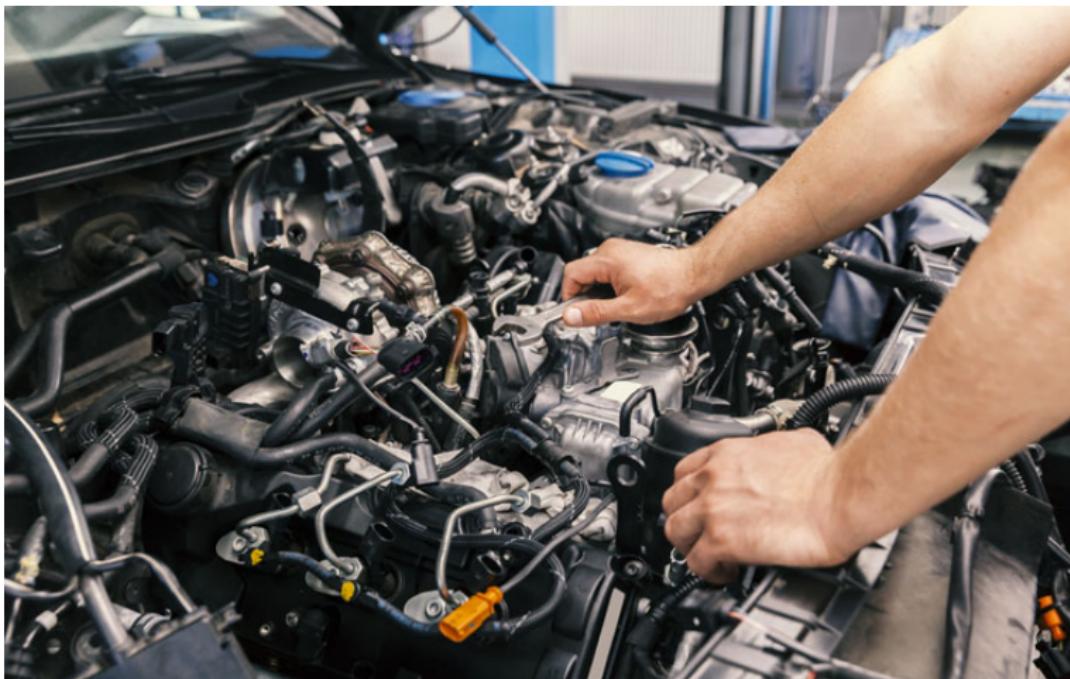


Usually





Not usually





Sometimes





Usually





Usually





Usually





Sometimes



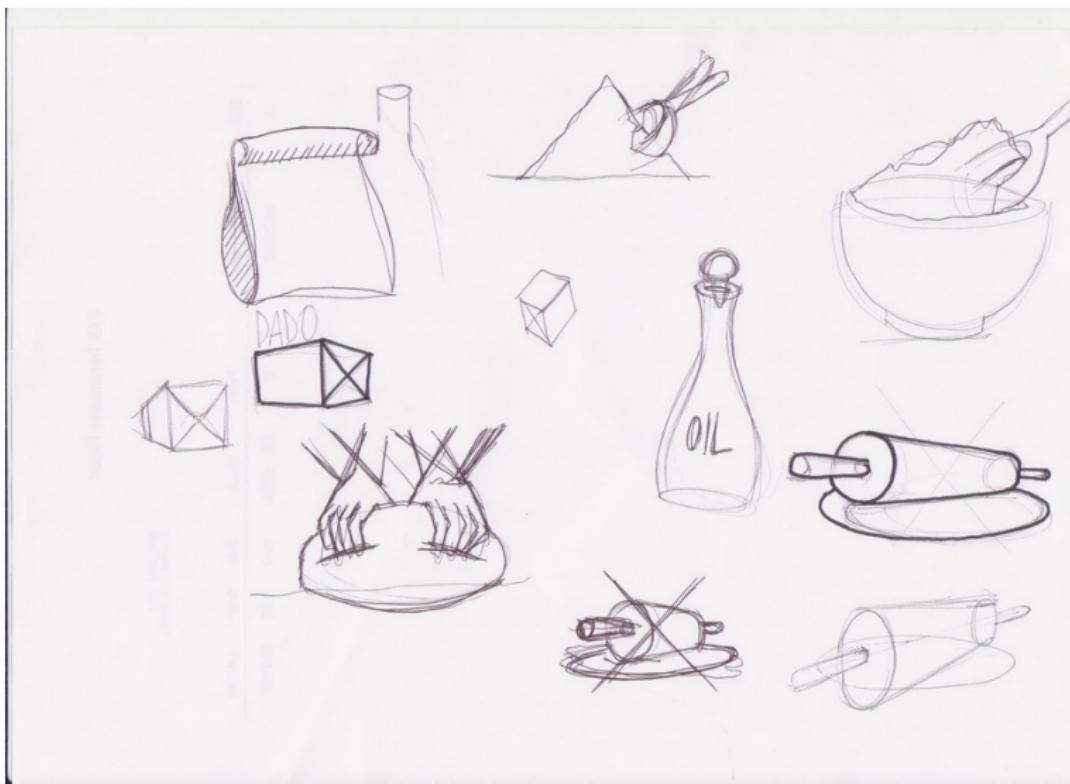


Sometimes





Better





They are all pizzas





They are all yeasts





The point

How much control?

- Do you want simply a result?
 - Be aware: even this is not always so easy in NoSQL era
- Do you want to be able to have intuitions?
- Do you want to control and analyze the whole process?



Variables

How can I control?

- Knowledge
- Time
- Access

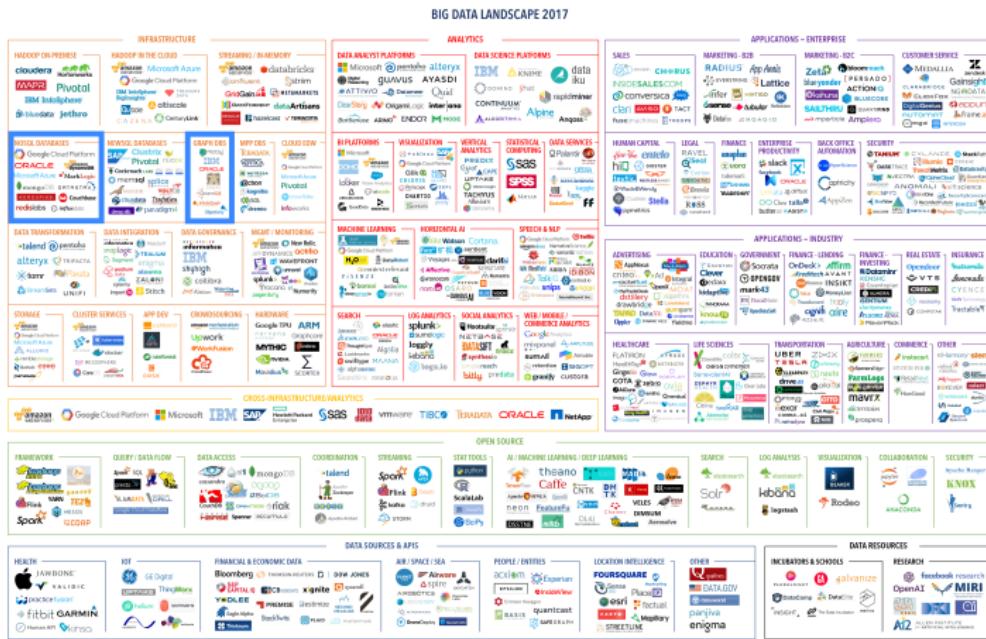


Big data is not just “Big”

Combining data from different sources

- **Machines**
 - real-time data, sensor data, trackers, streaming, etc.
- **People**
 - social media, personal documents, etc.
- **Organizations**
 - structured knowledge bases, data-warehouses, etc.

Technologies



Y2 - Last updated 5/2/2017

© Matt Turck (@mattturck), Jim Hao (@jimrhao), & FirstMark (@firstmarkcap) mattturck.com/bigdata2017

FIRSTMARK
EARLY STAGE VENTURE CAPITAL



Data management in Big Data

- How to ingest data
- How to store data
- How to ensure data quality
- How to identify the operations to be performed
- How to be efficient
- How to scale up data volume, variety, velocity, and access
- How to keep data secure



Data management in Big Data

- How to store data



Data management in Big Data

- How to store data
- How to scale up data volume, variety, velocity, and access



Big Data = Variety of Data Models

Identify the **different data models** that are used in the application

What is a data model?

- Structure
- Operations
- Constraints



Plan

1 Summary

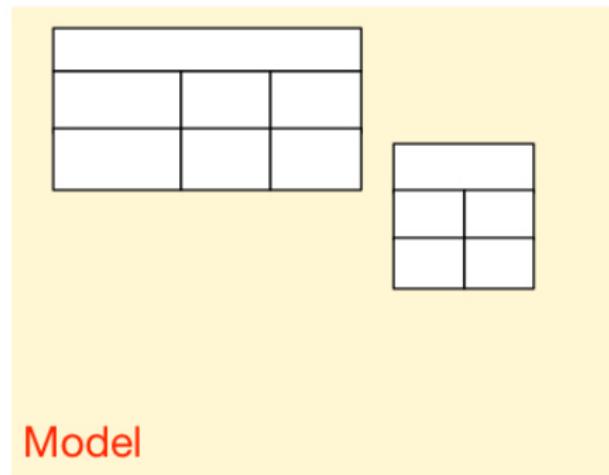
2 Relational Data Model

3 NoSQL

4 Data Modeling in NoSQL

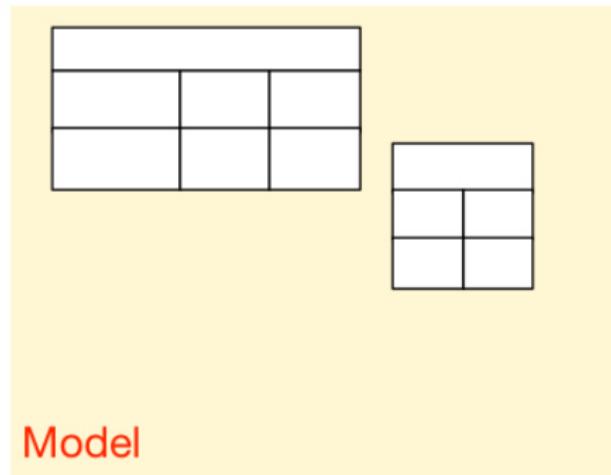


Relational Data model





Relational Data model



Structure

- Tables
- Columns



Relational Data model

Volges	Epinal	88	01
	...		
Moselle	Metz	51	01

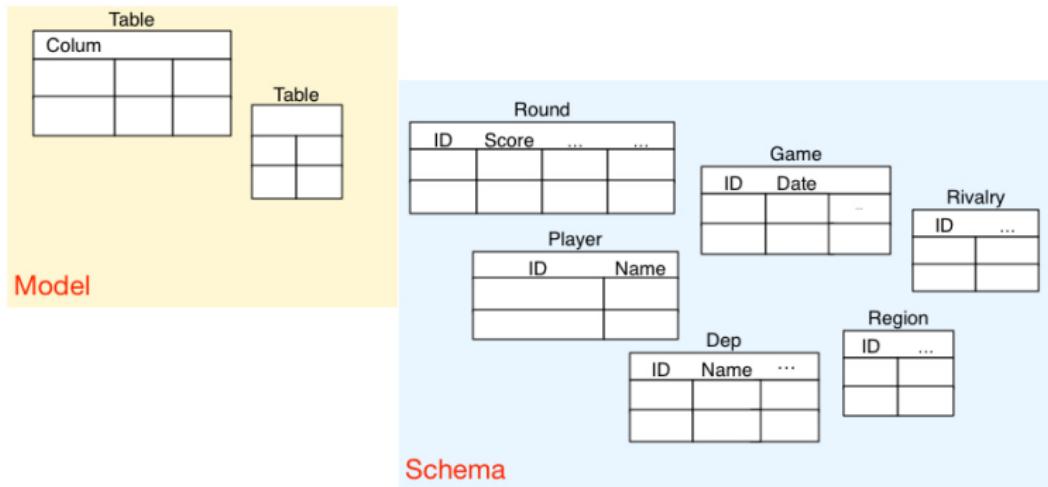
01	Grand Est
03	Occitanie

Structure

- Tuples
- Atomic values

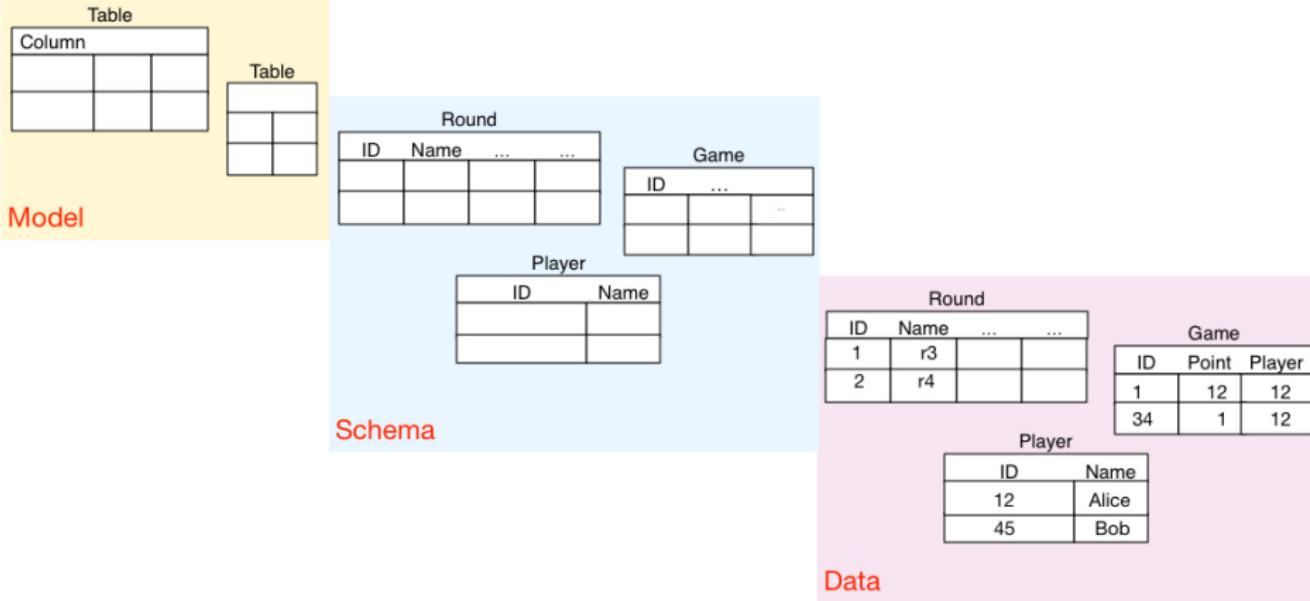


Relational Schema





Relational Data





Relational Data model

Relational modeling

It is possible to define tables in different ways but there are rules that govern the relational data modeling



Relational Data model

Game

PNick	PFName	PLName	...	PCity1	PCity1Reg	PCity2	PCity2Reg
Alice	Alice	Foo	...	Paris	Île-de-France	Lyon	Rhône-Alpes
Bob	Bob	Green	...	Nice	Aquitaine	Marseille	Aquitaine
Pt	Peter	Old	...	null	null	null	null

Data

All data in a big table?



Relational Data model

Game							
PNick	PFName	PLName	...	PCity1	PCity1Reg	PCity2	PCity2Reg
Alice	Alice	Foo	...	Paris	Île-de-France	Lyon	Rhône-Alpes
Bob	Bob	Green	...	Lyon	Rhône-Alpes	Marseille	Aquitaine
Pt	Peter	Old	...	null	null	null	null

Data

All data in a big table? What happens if..

- we add an address?



Relational Data model

Game

PNick	PFName	PLName	...	PCity1	PCity1Reg	PCity2	PCity2Reg
Alice	Alice	Foo	...	Paris	Île-de-France	Lyon	Rhône-Alpes
Bob	Bob	Green	...	Nice	Aquitaine	Marseille	Aquitaine
Pt	Peter	Old	...	null	null	null	null

Data

All data in a big table? What happens if..

- we add an address?
- some data is missing?



Relational Data model

Game							
PNick	PFName	PLName	...	PCity1	PCity1Reg	PCity2	PCity2Reg
Alice	Alice	Foo	...	Paris	Île-de-France	Lyon	Auvergne- Rhône- Alpes
Bob	Bob	Green	...	Lyon	Auvergne- Rhône- Alpes	Marseille	Nouvelle-Aquitaine
Pt	Peter	Old	...	null	null	null	null

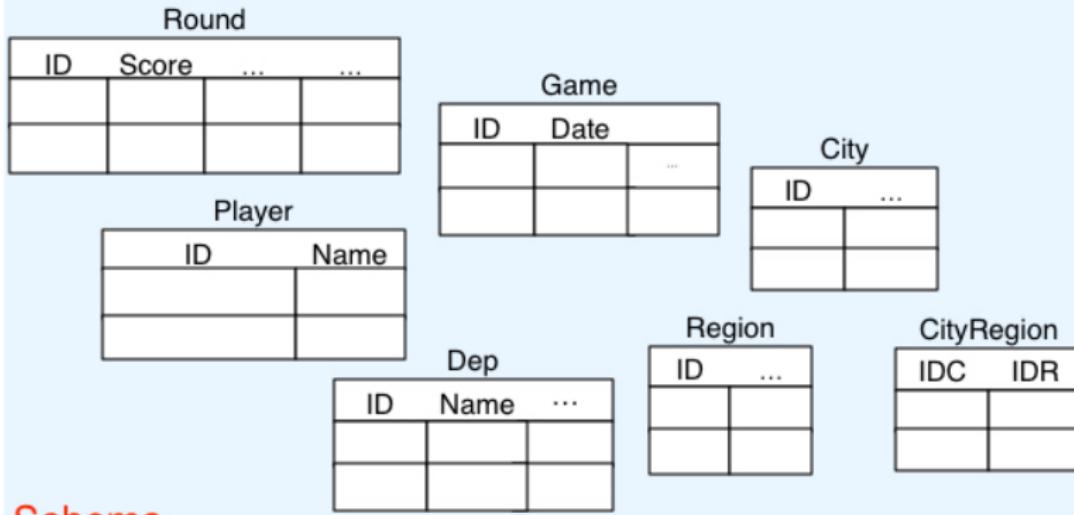
Data

All data in a big table? What happens if..

- we add an address?
- some data is missing?
- we update data?



Relational Data model





Relational Data model

Conceptual data modeling

- Each independent concept corresponds to a different table
- Constraints across tables that guarantee the data integrity and the links



Relational Data model

Candidate Key

uniquely identifies a tuple within a table. It can be composed by values belonging to one or multiple columns.

- No distinct tuples with the same values for its attributes
- It is minimal

Primary Key

in a DBMS is elected to be used as default tuple identifier for a table.

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Relational Data model

Foreign Key

is a value of a column (or a collection of values) in one table that uniquely identifies a row of another table or the same table

Round

ID	Score	WordSet	Timing	Game
rdf	23	we4	343	g34
...
fe4	56	4356re	654	g34
...
rde	45	3e4	321	g16

Game

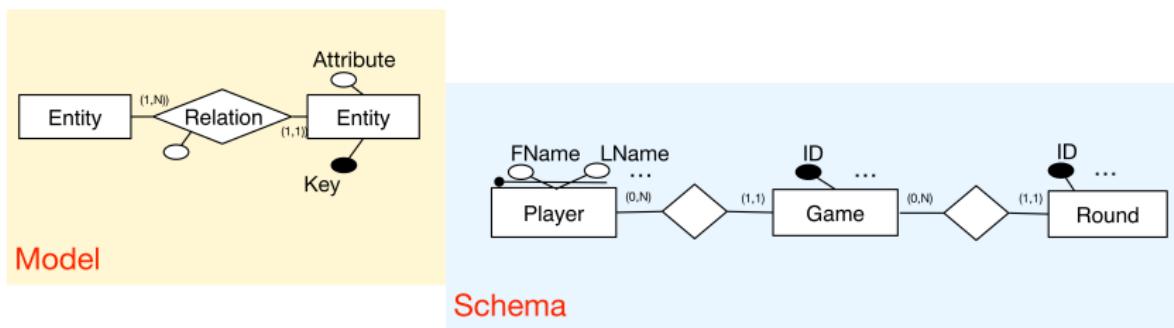
GID	PNick
g34	goodP
...
g16	theBest



Entity Relationship data model

Used to define a conceptual model

- Relationship
- Entity
- Attribute
- Key-attribute/Key-attributes
- Cardinalities





Relational Data model

Operations

Relational algebra operations (Codd [3] 1972) on tables (relations)

- **SELECT**: extracts tuples from a relation that satisfy a given restriction
- **PROJECT**: extracts specified attributes (columns) from a relation
- **PRODUCT**: builds the Cartesian product of two relations
- **UNION**: builds the set-theoretic union of two tables
- **INTERSECT**: builds the set-theoretic intersection of two tables
- **DIFFERENCE**: builds the set difference of two tables
- **JOIN**: connects two tables by their common attributes



Relational Data model

SQL Query Language

Structured Query Language (SQL) is the standard language for handling Relational data

- Data Definition Language
- Data Manipulation Language
- Data Control Language

Data Definition Language

Commands

CREATE TABLE (....)

ALTER TABLE (....)

DROP TABLE (....)

...

To create, modify, delete, etc. a table.



Example

Table creation

```
CREATE TABLE table_name (
    column1 datatype ,
    column2 datatype ,
    column3 datatype ,
    ...
);
```



Example

Table creation

```
CREATE TABLE Player
(
ID int,
name varchar(32)
);
```



Example

Table creation

```
CREATE TABLE Player
(
ID int,
name varchar(32)
);
```

Player	
ID	Name



Data Manipulation

SELECT (...)

DELETE (...)

UPDATE (...)

INSERT (...)

...



Example

Insert Data

```
INSERT INTO table_name (column1, column2, ...)
VALUES (value1, value2, value3, ...);
```



Example

Insert Data

```
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, value3, ...);
```

```
INSERT INTO Round (ID, Score, Wordset, Timing)  
VALUES (rdf, 45, we4, 343);
```

Round

ID	Score	WordSet	Timing
----	-------	---------	--------



Example

Insert Data

```
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, value3, ...);
```

```
INSERT INTO Round (ID, Score, Wordset, Timing)  
VALUES (rdf, 45, we4, 343);
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343

Round

ID	Score	WordSet	Timing
rdf	45	we4	343



Example

Data Manipulation

```
SELECT column1, column2, ...
FROM table_name;
```



Example

Data Manipulation

```
SELECT column1, column2, ...
FROM table_name;
```

```
SELECT *
FROM Round
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Example

Data Manipulation

```
SELECT ID , Score , Timing  
FROM Round
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Example

Data Manipulation

```
SELECT ID , Score , Timing  
FROM Round
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654

Round

ID	Score		Timing
rdf	45		343
rde	23		321
fe4	56		654



Example

Data Manipulation

```
SELECT ID, Score, Timing  
FROM Round  
WHERE Score > 40
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Example

Data Manipulation

```
SELECT ID, Score, Timing  
FROM Round  
WHERE Score > 40
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654

Round

ID	Score	WordSet	Timing
rdf	45		343
fe4	56		654



Example

Data Manipulation

```
SELECT ID , Timing  
FROM Round ,  
WHERE Score > 40
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Example

Data Manipulation

```
SELECT ID , Timing  
FROM Round ,  
WHERE Score > 40
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654

Round

ID			Timing
rdf			343
fe4			654



Example

Data Manipulation

```
SELECT ID , Timing  
      FROM Round ,  
 WHERE Score > 40 ,  
 AND Timing < 400
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654



Example

Data Manipulation

```
SELECT ID , Timing  
FROM Round ,  
WHERE Score > 40 ,  
AND Timing < 400
```

Round

ID	Score	WordSet	Timing
rdf	45	we4	343
rde	23	3e4	321
fe4	56	4356re	654

Round

ID			Timing
rdf			343



Example

Data Manipulation

```
SELECT ID , GID , PNick  
FROM Round , Game  
WHERE Round.Game = Game.GID
```

Round

ID	Score	WordSet	Timing	Game
rdf	23	we4	343	g34
...
fe4	56	4356re	654	g34
...
rde	45	3e4	321	g16

Game

GID	PNick
g34	goodP
...
g16	theBest



Example

Data Manipulation

```
SELECT ID , GID , PNick  
FROM Round , Game  
WHERE Round.Game = Game.GID
```

Round

ID	Score	WordSet	Timing	Game
rdf	23	we4	343	g34
...
fe4	56	4356re	654	g34
...
rde	45	3e4	321	g16

Game

GID	PNick
g34	goodP
...
g16	theBest

ID	GID	PNick
rdf	23	goodP
...
fe4	56	goodP
...
rde	45	theBest



Example

Data Manipulation

```
SELECT ID , GID , PNick  
FROM Round  
INNER JOIN Game ON Round.Game=Game.GID;
```

Round

ID	Score	WordSet	Timing	Game
rdf	23	we4	343	g34
...
fe4	56	4356re	654	g34
...
rde	45	3e4	321	g16

Game

GID	PNick
g34	goodP
...
g16	theBest



Example

Data Manipulation

```
SELECT ID , GID , PNick  
FROM Round  
INNER JOIN Game ON Round.Game=Game.GID;
```

Round

ID	Score	WordSet	Timing	Game
rdf	23	we4	343	g34
...
fe4	56	4356re	654	g34
...
rde	45	3e4	321	g16

Game

GID	PNick
g34	goodP
...
g16	theBest

ID	GID	PNick
rdf	23	goodP
...
fe4	56	goodP
...
rde	45	theBest



Examples

Data Control

```
CREATE TABLE Player
(
ID int,
name varchar(32) NOT NULL,
...
);
```

```
CREATE TABLE Player
(
ID int,
name varchar(32) NOT NULL,
...
PRIMARY KEY (ID)
);
```



And at last, what is a Database?

According to the classical literature a database is a

- Organized
- Shared
- Persistent
- Big

collection of data



And a Database Management System?

A database management system (DBMS) is a computer software system

Provides utilities for:

- Creating/managing large collections of data
- Accessing data collections
- Managing persistence
- Handling concurrent access
- Assuring reliability and privacy



Plan

- 1 Summary
- 2 Relational Data Model
- 3 NoSQL
- 4 Data Modeling in NoSQL



Why?

In the last thirty years relational databases have been the **default choice** for data storage

- Winner with respect to multiple alternatives during the years:
 - deductive databases in the 1980's
 - object databases in the 1990's
 - XML databases in the 2000's



Why?

In the last thirty years relational databases have been the **default choice** for data storage

- Winner with respect to multiple alternatives during the years:
 - deductive databases in the 1980's
 - object databases in the 1990's
 - XML databases in the 2000's
- A Data Scientist starting a new project:
 - which Relational database to use?
 - DB2, Oracle, SQLServer, MySQL, PostgreSQL, etc.
 - often no choice
 - Licenses
 - Triggers
 - Know-how

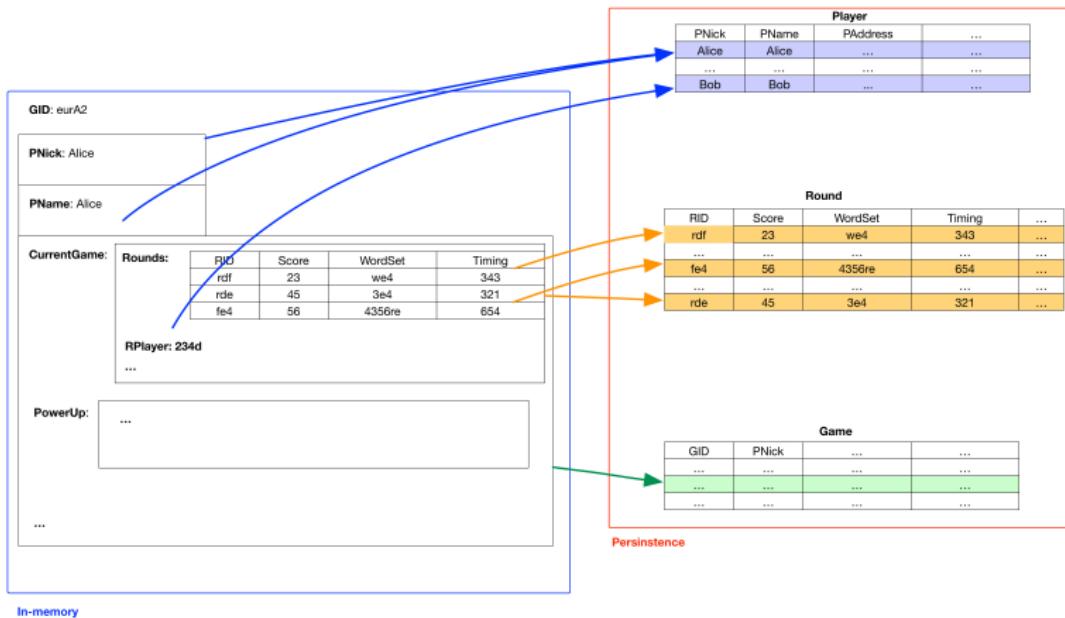


Provided Values

- Effective and efficient management of persistent data
- Concurrency control
- Data integration
- Standard data model
- Standard query language



Impedance Mismatch



Difference between the persistent data model and the in-memory data structures



Clusters

- New opportunities:
 - Data Volume
 - Commodity Hardware
- Relational databases were not designed to run on clusters
 - They do not perform well on clusters
- Increasing mismatch
 - Google Bigtable
 - Amazon DynamoDB



NoSQL

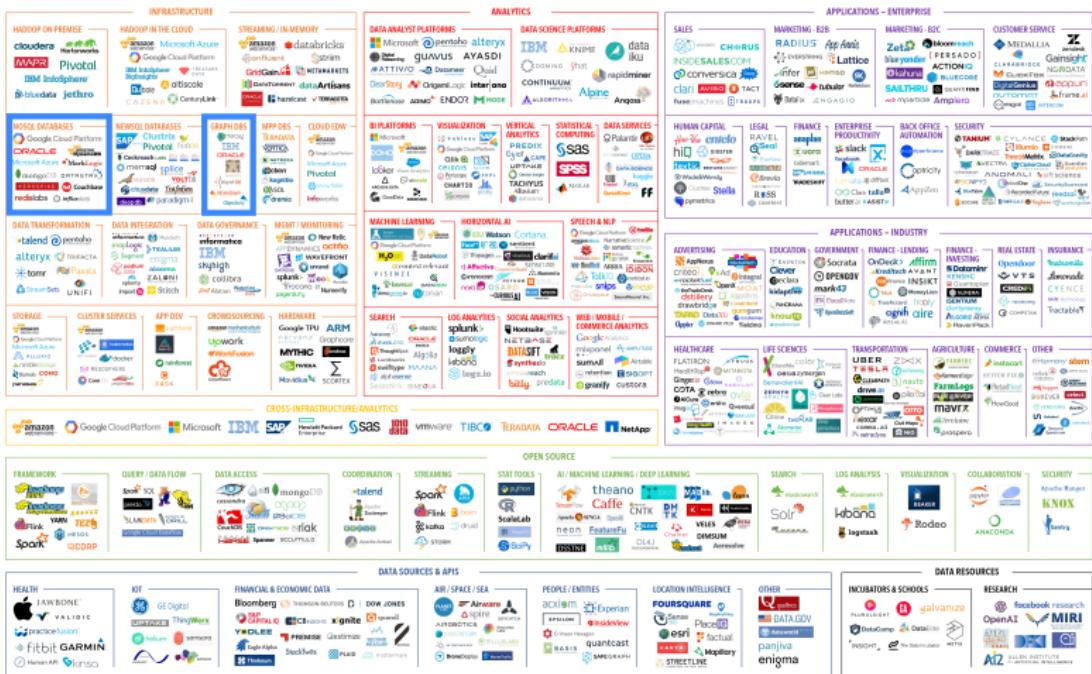


- Term appeared in the late 90s
 - Open source relational database [Strozzi NoSQL]
 - Tables as ASCII files, without SQL
- Current interpretation:
 - 11 June, 2009: San Francisco meet-up
 - Hashtag chosen NoSQL
 - Main Features:
 - Not using the Relational Model and SQL
 - Open source projects (mostly)
 - Running on clusters
 - Schema-less

NoSQL



NoSQL



M2 – Last updated 5/3/2012

© Matt Turck (@mattturck), Jim Hao (@jimrhao), & FirstMark (@firstmarkcap) mattturck.com/bigdata2017

FIRSTMARK
EARLY STAGE VENTURE CAPITAL

NoSQL

NoSQL Databases



amazon
DynamoDB



Google Cloud Platform

ORACLE

Microsoft Azure



mongoDB



MarkLogic



DATASTAX

AEROSPIKE



Couchbase



SequoiaDB

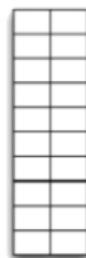
redislabs



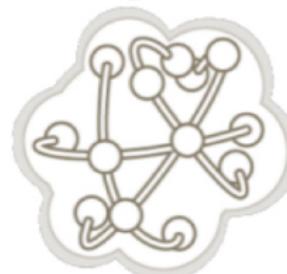
influxdata



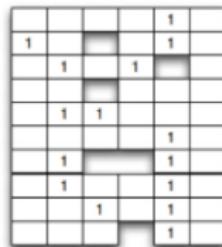
NoSQL



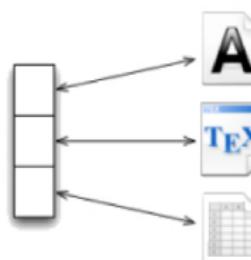
Key-Value



Graph



Column-Family



Document



Data Models For NoSQL Systems

- Document
 - MongoDB
- Key-value
 - Redis
- Column-family
 - Cassandra
- Graph
 - Neo4j



Pay Attention

Many possible **data representations** are possible

There is **no theory** that can state which data representation is the best of all for a given data store



Pay Attention

collection Player

<i>id</i>	<i>document</i>
8*mary	{ _id:"mary", username:"mary", firstName:"Mary", games[1]: { game:"Game:2345", opponent:"Player:rick" }, games[2]: { game:"Game:7425", opponent:"Player:ann" } }

collection Player

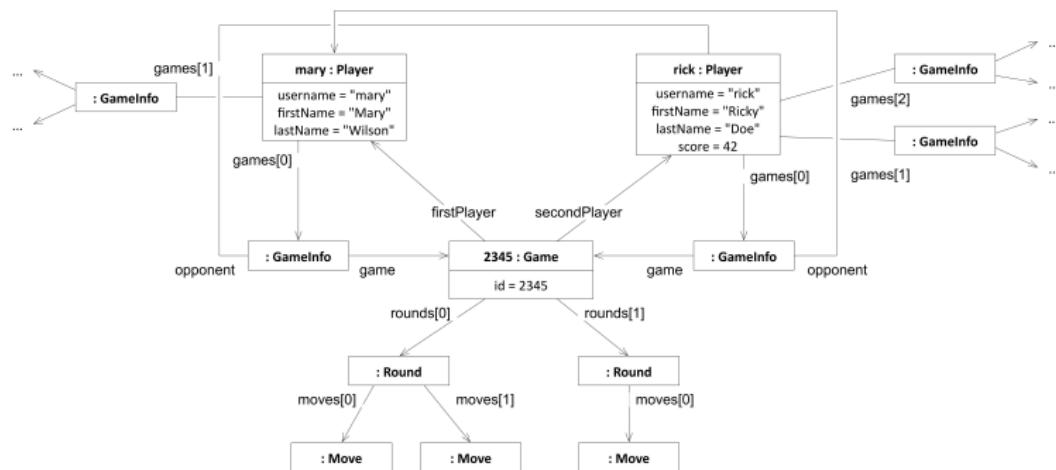
<i>id</i>	<i>document</i>
7*mary	{ _id:"mary", username:"mary", firstName:"Mary", games: [{ game:"Game:2345", opponent:"Player:rick"}, { game:"Game:7425", opponent:"Player:ann"}] }

Figure: Two possible data representations in MongoDB

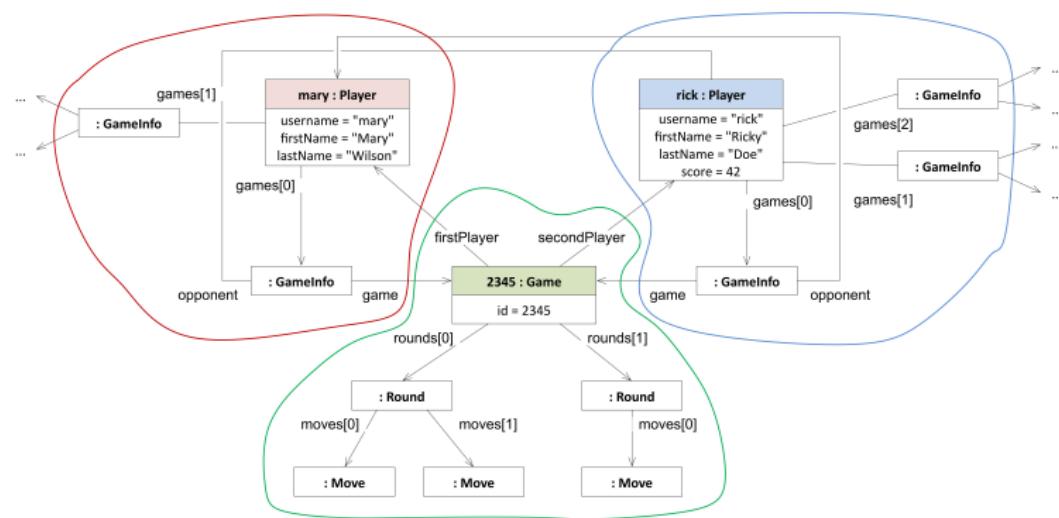
Aggregates

- Data as units that have a complex structure
 - More structure than just a set of tuples
 - A complex record with fields inside
 - A record with a composite access structure
- Aggregates in Domain Driven Design [4]:
 - A collection of related objects that we treat as a unit
 - A unit for data manipulation and management of consistency
- Advantages of aggregates:
 - easier for programmer to work with
 - easier for database systems to handle operations on a cluster

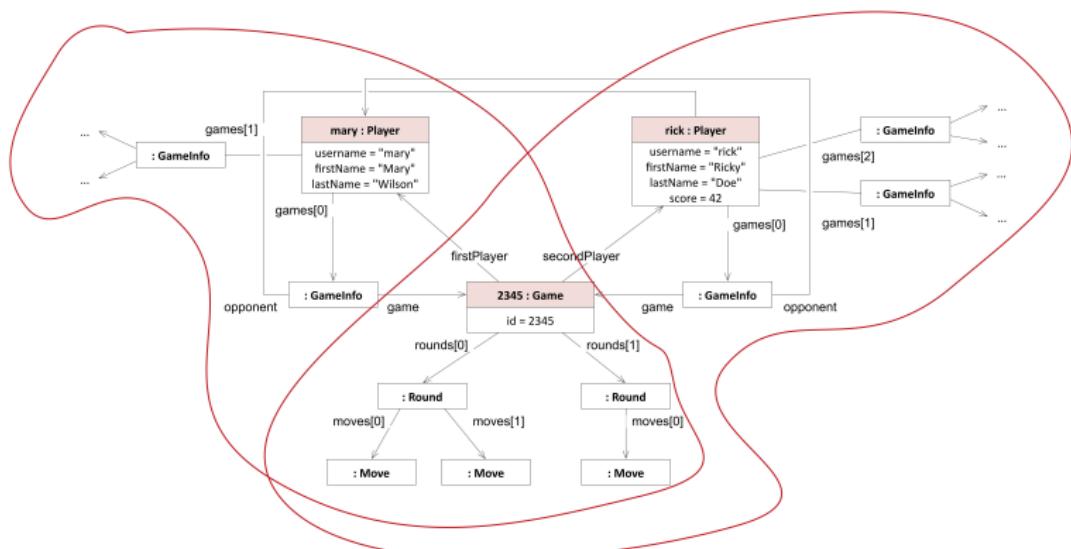
Aggregates [1]



Aggregates



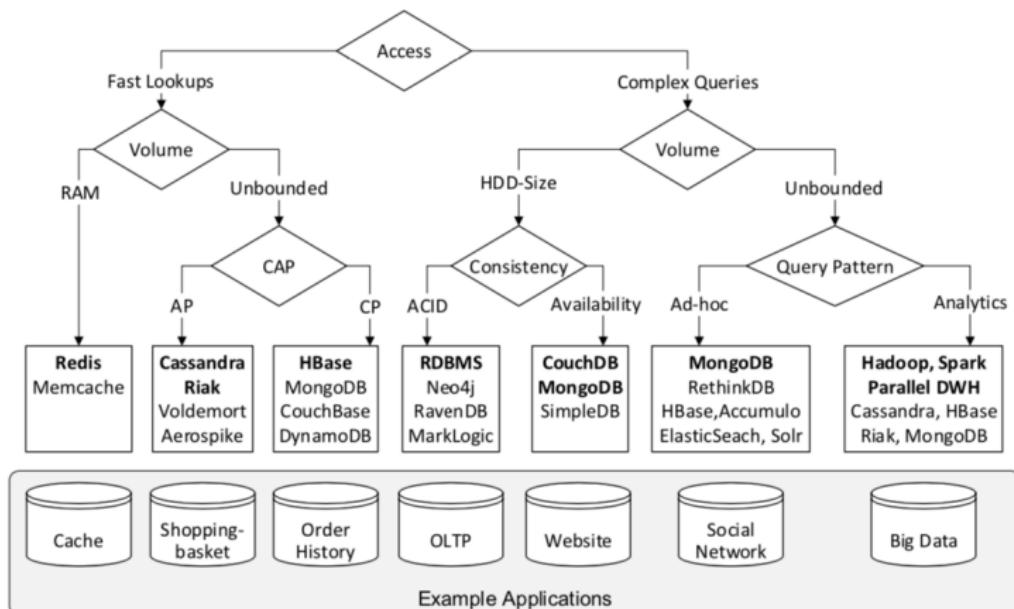
Aggregates



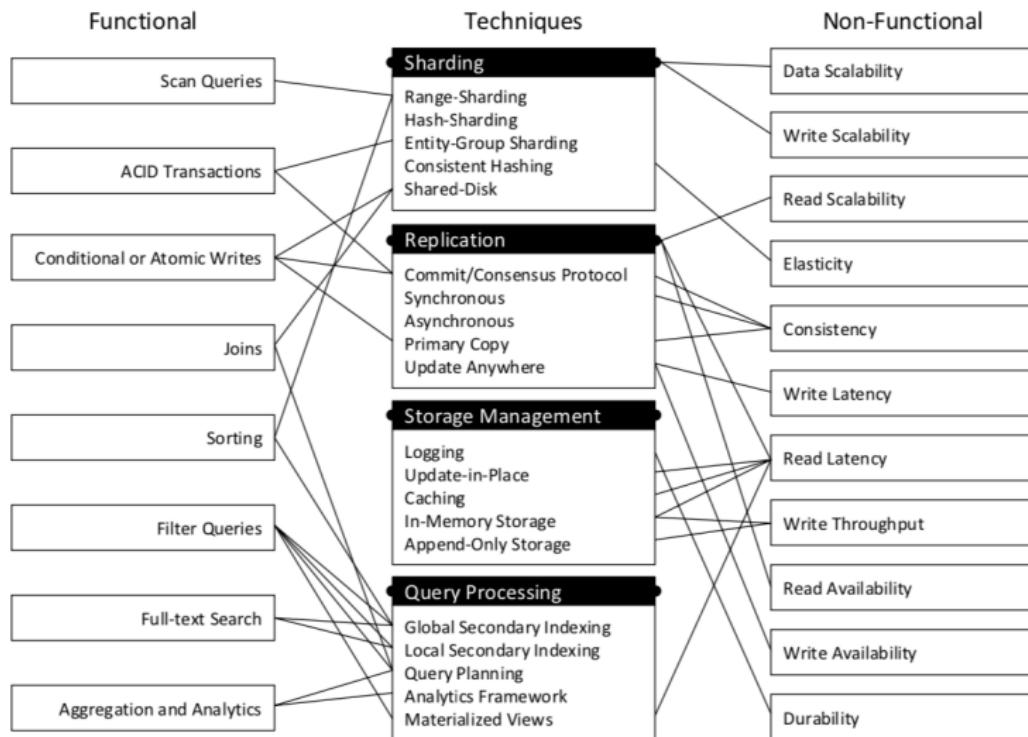
Design Strategy

- **No universal strategy** for designing aggregates boundaries
- The design depends on **data manipulation**
 - Access on a single round
 - Access on a player with all his games and rounds
- Context specific:
 - Some applications will run on a data representation
 - Some others on another data representation
- Focus on the **Data Access Unit**
 - very good for clusters
 - not easy to change the data interaction strategy

Not easy to classify and multiple options possible[6]



Not easy to classify and multiple options possible[6]





Not easy to classify and multiple options possible[6]

Query Capabilities



Not easy to classify and multiple options possible[6]

Table 1 Querying capabilities

NoSQL Data Stores	Querying Map Reduce	Query			Other API	Other features	License
		REST	SQL	Other			
Key-value stores	Redis http://redis.io	No	Third-party APIs	Does not provide SQL-like querying	CLI and API in several languages	Server-side scripting support using Lua	Open source: BSD (Berkeley Software Distribution)
	Memcached http://memcached.org	No	Third-party APIs	Does not provide SQL-like querying	CLI and API in several languages. Binary and ASCII protocols for custom client development	No server-side scripting support	Open source: BSD 3-clause license
	BerkeleyDB http://www.oracle.com/us/products/database/berkeley-db/www/index.html	No	Yes	SQLite	CLI and API in several languages	No secondary indices, no server-side-scripting support	Closed source: Oracle Sleepycat license
	Voldemort http://www.project-voldemort.com/voldemort	Yes	Under development	No	Clients for several languages		Open source: Apache 2.0 license
Column family stores	Riak http://basho.com/riak	Yes	Yes	Rank search, secondary indices	CLI and API in several languages	Provides filtering through key filters. Configurable secondary indexing. Provides full search capabilities. Provides server-side scripting	Open source: Apache 2.0 license
	Cassandra http://cassandra.apache.org	Yes	Third party APIs	Cassandra query language	CLI and API in several languages. Supports Thrift interface	Secondary indexing mechanisms include column families, super-columns, column ranges	Open source: Apache 2.0 license
	HBase http://hbase.apache.org	Yes	Yes	No, could be used with Hive	Java/Any Writer	Server-side scripting support. Several secondary indexing mechanisms	Open source: Apache 2.0 license
	DynamoDB (Amazon service) http://aws.amazon.com/dynamodb	Amazon Elastic MapReduce	Yes	Proprietary	API in several languages	Provides secondary indexing based on attributes other than primary keys	Closed source: Pricing as per-use basis
Document stores	Amazon SimpleDB (Amazon service) http://aws.amazon.com/simpledb	No	Yes	Amazon proprietary	Amazon proprietary API	Automatic indexing for all columns	Closed source: Pricing as per-use basis
	MongoDB http://www.mongodb.org	Yes	Yes	Proprietary	CLI and API in several languages	Server-side scripting and secondary indexing support. A powerful aggregation framework	Open source: Free GNU AGPL v3.0 license
	CouchDB http://couchdb.apache.org	Yes	Yes	SQL like SQL, under development	API in several languages	Server-side scripting and secondary indexing support	Open source: Apache 2.0 license
	Couchbase Server http://www.couchbase.com	Yes	Yes	No	Memcached API + protocol changes and ASCII in several languages	Server-side scripting and secondary indexing support	Open source: Free Community Edition. Paid Enterprise Edition
Graph databases	Neo4j http://www.neo4j.org	No	Yes	Cypher, Gremlin and SparQL	CLI and API in several languages	Server-side scripting and secondary indexing support	Open source license: NTCL + (AGPLv3)
	HyperGraphDB http://www.hypergraphdb.org/	No	Yes	SQL like querying	Currently has Java API. Could be used with Scala	Provides a search engine and Sesco scripting	Open source license: GNU IDE
	AllegroGraph http://www.franz.com/graph/allegrograph	No	Yes	SparQL and Prolog	API in several languages	Support for Solr indexing and search	Closed source: free developer and enterprise versions
	Voltdb http://voltmdb.com	No	Yes	SQL	CLI and API in several languages. JDBC support	StoredProcedure are written in Java. Tables cannot join with themselves, and all joined tables must be partitioned over the same value	Open source: AGPL v3.0 license. Commercial enterprise edition
NewSQL	Spanner	Yes	NA	SQL like language	NA	Tables are partitioned into hierarchies, which describe locality relationship between tables	Google internal use only
	Choraix http://www.choraix.com/	No	No	SQL	Wire protocol compatible with MySQL		Closed source. Available as a service in the AWS marketplace. Available as an appliance, and as standalone software
	NuoDB http://www.nuodb.com/	No	No	SQL	CLI and drivers for most common data access APIs (JDBC, ODBC, ADO.NET). Also provides a C++ API	No support for stored procedures	Closed source. Pro and Developers editions. Available as a service in the AWS marketplace



Not easy to classify and multiple options possible[6]

Partition Replication and Consistency



Not easy to classify and multiple options possible[6]

Table 2 Partitioning, replication, consistency, and concurrency control capabilities

		Partitioning	Replication	Consistency	Concurrency control
Key-value stores	Redis	Not available (planned for Redis Cluster release). It can be Master-slave, asynchronous replication, implemented by a client or a proxy.	No replication.	Eventual consistency. Strong consistency if slave replicas are solely for failover.	Application can implement optimistic (using the WATCH command) or pessimistic concurrency control.
	Memcached	Clients' responsibility. Most clients support consistent hashing.	No replication. Replicated can be added to memcached for replication.	Strong consistency (single instance).	Application can implement optimistic (using CAS with version stamps) or pessimistic concurrency control.
	BerkeleyDB	Key-range partitioning and custom partitioning functions. Not supported by the C# and Java APIs at this time.	Master-slave.	Configurable.	Readers-writer locks.
	Voldemort	Consistent hashing.	Masterless, asynchronous replication. Replicas are located on the first <i>R</i> nodes moving over the partitioning ring in a clockwise direction.	Configurable, based on quorum read and write requests.	MVCC with vector clock.
	Riak	Consistent hashing.	Masterless, asynchronous replication. The built-in functions determine how replicas distribute the data evenly.	Configurable, based on quorum read and write requests.	MVCC with vector clock.
Column family stores	Cassandra	Consistent hashing and range partitioning (known as order preserving hashing in Cassandra terminology) is not recommended due to the possibility of hot spots and load balancing issues.	Masterless, asynchronous replication. Replicas are replicated. Replicas are placed on the next <i>R</i> nodes along the ring, or, replica 2 is placed on the first node along the ring that belongs to another data centre, with the remaining replicas on the nodes along the ring in the same rack as the first.	Configurable, based on quorum read and write requests.	Client-provided timestamps are used to determine the most recent update to a column. The latent timestamp always wins and eventually persists.
	HBase	Range partitioning.	Master-slave or multi-master, asynchronous replication. Does not support horizontal load-balancing (a row is served by exactly one server). Replicas are used only for failover.	Strong consistency.	MVCC.
	DynamoDB	Consistent hashing.	Three-way replication across multiple zones in a region. Synchronous replication.	Configurable.	Application can implement optimistic (using incrementing version numbers) or pessimistic concurrency control.
	Amazon SimpleDB	Partitioning is achieved in the DB design stage by normally adding additional domains (tables). Cannot query across domains.	Replicas within a chosen region.	Configurable.	Application can implement optimistic concurrency control by specifying a version number (or a timestamp) attribute and by performing a conditional put/delete based on the attribute value.
Document Stores	MongoDB	Range partitioning based on a shard key (one or more fields that exist in every document in the collection). In addition, hashed shard keys can be used to partition data.	Master-slave, asynchronous replication.	Configurable. Two methods to achieve strong consistency: set connection to read only from primary, or, set write concern parameter to "Replica Acknowledged".	Readers-writer locks.
	CouchDB	Consistent hashing.	Multi-master, asynchronous replication. Designed for off-line operation. Multiple replicas can maintain their own copies of the same data and synchronize them at a later time.	Eventual consistency.	MVCC. In case of conflicts, the winning revision is chosen, but the losing revision is saved as a previous version.
	Cassandra Server	A hashing function determines to which bucket a document belongs. Next, a table is consulted to look up the server that hosts that bucket.	Multi-master.	Within a cluster: strong consistency. Across clusters: eventual consistency.	Application can implement optimistic (using CAS) or pessimistic concurrency control.
Graph databases	Neo4J	No partitioning (cache sharding only).	Master-slave, but can handle write requests on all server nodes. Write requests to slaves must synchronously propagate to master.	Eventual consistency.	Write locks are acquired on nodes and relationships until committed.



Not easy to classify and multiple options possible[6]

Security

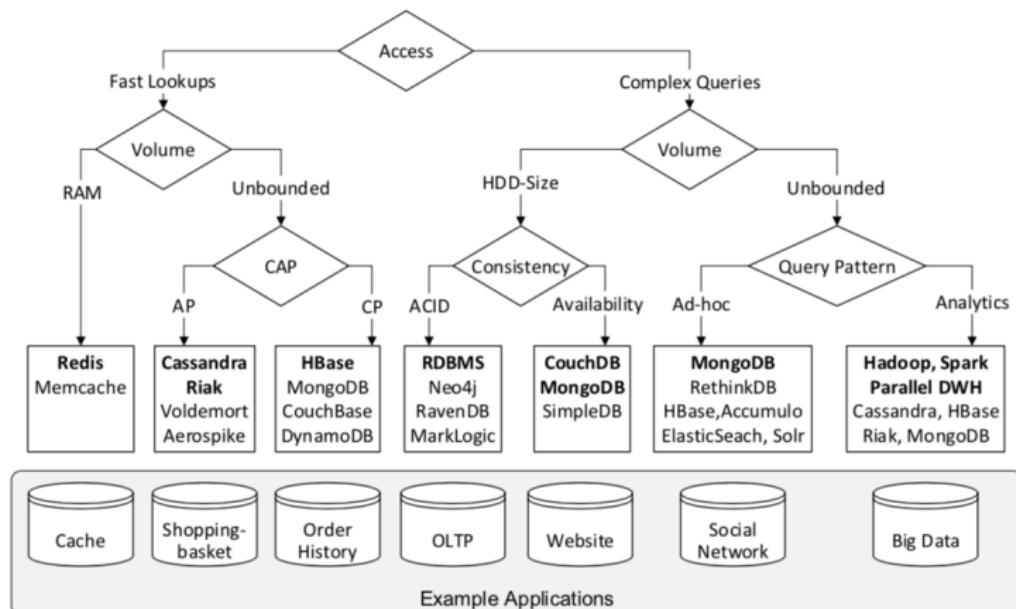


Not easy to classify and multiple options possible[6]

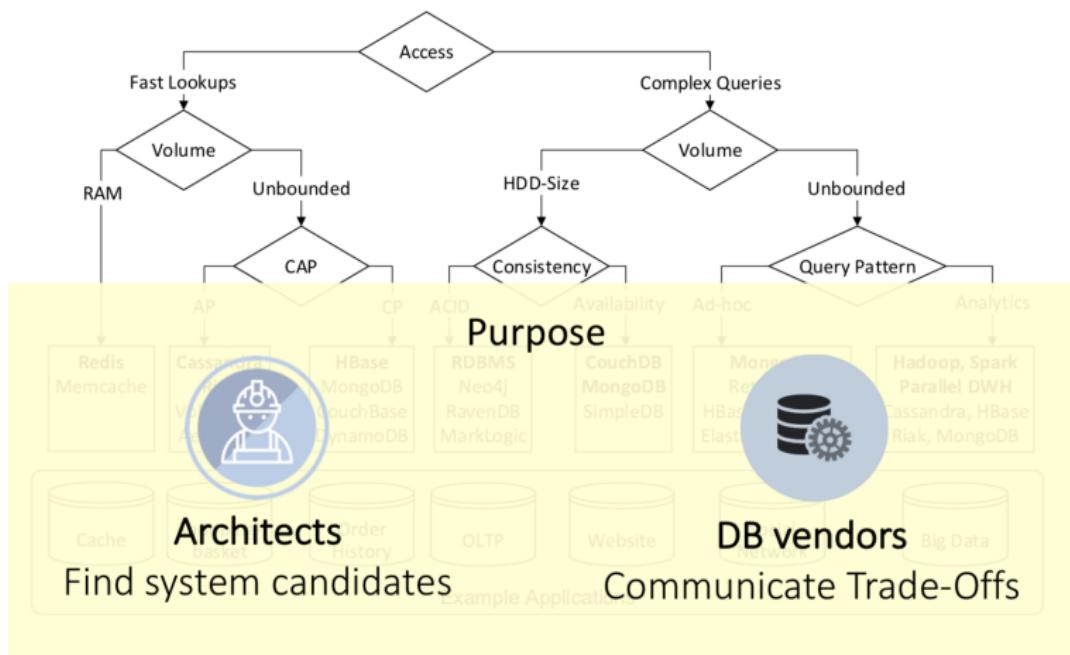
Table 3 Security features

NoSQL Data Stores		Encryption Data at Rest	Client/Server	Server/Server	Authentication	Authorization	Auditing
Key-Value stores	Redis	No	No	No	Admin password sent in clear text for admin functions. Data access does not support authentication.	No	No
	Memcached	NA, Memcache does store data on disk	No	No	Binary protocol supports Simple Authentication and No Security Layer (SASL) authentication	No	No
	BerkeleyDB	Yes, database needs to be created using encryption	NA, embedded data store	No	No	No	No
	Voldemort	Possibly if BerkeleyDB is used as the storage engine	No	No	No	No	No
	Riak	No REST interface supports HTTPS. Binary protocol is not encrypted	Multiple data-center replication can be done over HTTPS	No	No	No	No
Column Family Stores	Cassandra	Enterprise Edition only. Commit log is not encrypted	Yes, SSL based	Yes, configurable: all server-to-server communication, only between datacenters or between servers in the same rack	Yes, store credentials in a system table. Possible to provide pluggable implementations	Yes, similar to the SQL GRANT/REVOKE approach. Possible to provide pluggable implementations	Enterprise Edition only. Based on log4j framework. Logging categories include ADMIN, ALL, AUTH, DML, DDL, DCL, and QUERY. Possibility to disable logging for specific keyspaces
	HBase	No, planned for future release	Yes	Communication of Hbase nodes with the HDFS and Zookeeper clusters can be Kerberos, REST API, or a HTTP gateway, which authenticates with its data center as one single user, and executes all operations on his/her behalf	Yes, RPC API based on SASL, supporting Kerberos, REST API, or a HTTP gateway, which authenticates with its data center as one single user, and executes all operations on his/her behalf	Yes, permissions include read, write, create and admin. Granularity of table, column family, or column	No, planned for future release
	Amazon DynamoDB	No	Yes, HTTPS	NA	Integration with Identity and Access Management (IAM) service. The requests need to be signed using HMAC-SHA256	Allow the creation of policies that associate users and operations on domains. Possible to define policies for temporary access	Integrates with Amazon Cloud Watch service. Access information about latencies for operations, amount of data stored, and requests throughput
Document Stores	Amazon SimpleDB	See DynamoDB					No
	MongoDB	No, a third-party partner (Gazzang) provides an encryption plug-in	Yes, SSL-based	Yes	Yes, store credentials in a system collection. REST interface does not support authentication. Enterprise Edition supports Kerberos	Yes, permissions include read, write, update, dbAdmin, and userAdmin. Granularity of collection	No
	CouchDB	NA	Yes, SSL-based	Possible using HTTPS connections	Yes, HTTP authentication using cookies or BASIC method. OAuth supported	Three levels of users: server admin, database admin, and database member. Complex authorization can be done in validation functions	No
	Couchbase Server	No	No	No, planned for future release	Yes, SASL authentication – each bucket is differentiated by its name and password. REST API for administrative function uses HTTP BASIC authentication	No	No
Graph databases	Nos4J	No	Yes, SSL-based	No	No, developers can create a SecurityRule and register with the server	No	No

Not easy to classify and multiple options possible[6]

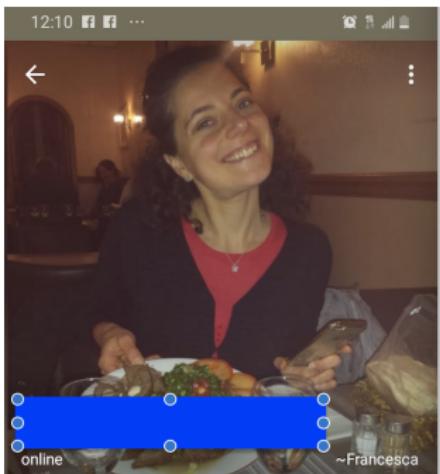


Not easy to classify and multiple options possible[6]





In practice? How do they work?



12:10 f f ... 4G 09:08 8%

Telekom.de 4G

Impostazioni Modifica profilo

Inserisci il tuo nome e (facoltativo) un'immagine per il tuo profilo

Modifica

Francesca

NUMERO DI TELEFONO

INFO

Media visibility Big Data

Francesca Bugiotti

86/155

Media

68 >

online ~Francesca

Mute notifications

Custom notifications

lion emoji 😊 sun emoji ☀️ "E quindi uscimmo a riveder le st... >



Plan

- 1 Summary
- 2 Relational Data Model
- 3 NoSQL
- 4 Data Modeling in NoSQL



Document Database Word





Document Database

The database is a collection of documents

- The database stores and retrieves **documents**
- A key-value data-store where the value is an **examinable** document
- Documents:
 - can be XML, JSON, BSON, and so on
 - are self-describing, hierarchical tree data structures
 - can consist of scalar values, collections, and maps
 - are structurally similar but not identical to each other

Operations:

- **get, put, delete** a document given its ID
- **query** the content of the document



Example

```
{  
  "firstname": "Pramod",  
  "gamesPlayed": [ "GameStar",  
                  "GamePuzzle",  
                  "GameCities"],  
  "addresses": [  
    { "state": "France",  
      "city": "Paris",  
      "Department": "Essonne"  
    },  
    { "state": "France",  
      "city": "Lille",  
      "type": "Nord" }  
  ],  
}
```



Document Databases

- Strongly aggregate oriented
 - A lot of aggregates
 - Each aggregate has a key
- Data Model:
 - A set of (key,document) pairs
 - Document: an aggregate instance
- It is possible to query/navigate the aggregate
- Access to an aggregate:
 - queries based on the **fields of the aggregate**



Example

collection Player

<i>id</i>	<i>document</i>
8*mary	{ _id:"mary", username:"mary", firstName:"Mary", games[1]: { game:"Game:2345", opponent:"Player:rick" }, games[2]: { game:"Game:7425", opponent:"Player:ann" } }

collection Player

<i>id</i>	<i>document</i>
7*mary	{ _id:"mary", username:"mary", firstName:"Mary", games: [{ game:"Game:2345", opponent:"Player:rick"}, { game:"Game:7425", opponent:"Player:ann"}] }

Figure: Two possible data representations in MongoDB



MongoDB



Database structure

- Each MongoDB instance has multiple **databases**
 - Similar to a database schema in a RDBMS
- Each database can have multiple **collections**
 - Similar to a table in a RDBMS
- When we store a document, we have to choose which database and collection this document belongs to
`db.collection.insert(document)`



MongoDB

Document Structure

- The **schema** of the data in a collection can differ across documents
 - It is possible to say that a collection is a set of **similar** documents
- In documents no empty attributes
 - Empty values correspond to not-included attributes
- New attributes can be created without the need to define them or to change the existing documents
- It is possible to include special attributes that reference other documents



Example

```
{  
    "_ID": 2330,  
    "firstname": "Alice",  
    "lastname": "Foo",  
    "gamesPlayed": [ "GameStar",  
                     "GamePuzzle",  
                     "GameCities" ],  
    "addresses": [  
        { "state": "France",  
          "city": "Paris",  
          "Department": "Essonne" },  
        { "state": "France",  
          "city": "Lille",  
          "type": "Nord" }  
    ],  
    "birthday": "25 August"  
}
```

```
{  
    "_ID": 2332,  
    "firstname": "Pramod",  
    "gamesPlayed": [ "GameStar",  
                     "GamePuzzle",  
                     "GameCities" ],  
    "addresses": [  
        { "state": "France",  
          "city": "Paris",  
          "Department": "Essonne" },  
        { "state": "France",  
          "city": "Lille",  
          "type": "Nord" }  
    ]  
}
```



Example

```
{ _ID: 2330,  
  "firstname": "Alice",  
  "lastname": "Foo"  
  "gamesPlayed": [ "GameStar",  
                   "GamePuzzle",  
                   "GameCities"],  
  "games": [ 4657,  
            2345,  
            56984],  
  "birthday": "25 August"  
}
```

```
{ _ID: 2345,  
  "gameName": "GameStar",  
  "rounds": [ 4564, 45696],  
  "globalScore": 45565},  
  "opponent": 445,  
  "globalTime": "55769"  
}
```



Query features

In General

You can query the data inside the document without having to retrieve the whole document by its key and then introspect the document

Different document DBMS provide different query features



MongoDB query language

Expressed via JSON with simple constructs

`find()` operation

```
db.collection.find( <query filter>, <projection>
).<cursor modifier>
```

```
/* in games
{
  "gameId": "99",
  "playerId": "1234fb",
  "gameDate": "2014-04-25",
  "rounds": [
    { "round": { "id": 27, "score": "45"}, "time": 32.45 },
    { "round": { "id": 55, "score": "75"}, "time": 41.33 }
  ],
}
db.games.find()
```

Matches every document in the collection `games`



MongoDB query language

Expressed via JSON with simple constructs

`find()` operation

```
db.collection.find( <query filter>, <projection>
).<cursor modifier>
```

```
/* in games
{
  "gameId": "99",
  "playerId": "1234fb",
  "gameDate": "2014-04-25",
  "rounds": [
    { "round": { "id": 27, "score": "45"}, "time": 32.45 },
    { "round": { "id": 55, "score": "75"}, "time": 41.33 }
  ],
}
*/
db.games.find({ "playerId": "1234fb" })
```

Finds all the documents where the value for playerId is 234fb



MongoDB query language

Expressed via JSON with simple constructs

`find()` operation

```
db.collection.find( <query filter>, <projection>
).<cursor modifier>
```

```
/* in games
{
  "gameId": "99",
  "playerId": "1234fb",
  "gameDate": "2014-04-25",
  "rounds": [
    { "round": { "id": 27, "score": "45"}, "time": 32.45 },
    { "round": { "id": 55, "score": "75"}, "time": 41.33 }
  ],
} */
db.games.find({ 'playerId': '1234fb' }, { 'gameId': 1, 'gameDate': 1 })
```

The query returns only the attributes `_id` (returned by default), `gameID`, and `gameDate`



MongoDB query language

Expressed via JSON with simple constructs

`find()` operation

```
db.collection.find( <query filter>, <projection>
).<cursor modifier>
```

```
/* in games
{
  "gameId": "99",
  "playerId": "1234fb",
  "gameDate": "2014-04-25",
  "rounds": [
    { "round": { "id": 27, "score": "45"}, "time": 32.45 },
    { "round": { "id": 55, "score": "75"}, "time": 41.33 }
  ],
}
*/
db.games.find({ 'rounds.time': 32.45 }, { playerId: 1, rounds: { $ : 1 } })
```

We can query also according the value of some attribute contained in an array.



MongoDB query language

Expressed via JSON with simple constructs

`find()` operation

```
db.collection.find( <query filter>, <projection>
).<cursor modifier>
```

```
/* in games
{
  "gameId": "99",
  "playerId": "1234fb",
  "gameDate": "2014-04-25",
  "rounds": [
    { "round": { "id": 27, "score": "45"}, "time": 32.45 },
    { "round": { "id": 55, "score": "75"}, "time": 41.33 }
  ],
}
*/
db.games.find({ "rounds": { "$elemMatch": { "round.score": { "$gte": 45 }}}})
```

We can query also embedded documents:

- we can impose a condition on the score value



MongoDB[5]



Column-family databases



Google
BigTable



HYPERTABLE^{INC}

Column-family databases I

Also called **Extensible Record Stores**, they come after the Google's success with BigTable [2]

Their basic data model consists of **rows** and **columns**

- **Rows** are analogues to documents
 - **Variable number** of columns (called also attributes, fields, qualifiers, etc.) having a name
 - Any type for the columns
 - Rows are grouped into **collections** (called also tables)
 - Values associated with timestamps
- **Columns** are grouped together and form **column groups** (super columns, column families, etc.)
 - a column group is often a namespace for a column
 - columns with the same name in different column groups



Column-Family Databases

- Strongly aggregate oriented
 - A lot of aggregates
 - Each aggregate has a key
- Data Model:
 - A two level map structure
 - A set of (row-key, aggregate) pairs
 - Each aggregate is a group of pairs (column-key, value)
- Structure of the aggregate is visible
- Columns can be organized in families
 - Data usually accessed together
- Access to an aggregate:
 - a row as a whole
 - part of a column



Example

[Item per object in DynamoDB]

table	_id	username	firstName	lastName	games
Player	mary1994	mary1994	Mary	Wilson	{ ... }

[Cell per object in Cassandra]

table	<i>id</i>	<i>value</i>
Player	mary1994	"username":"mary1994", "firstName":"Mary", ...



Column-family databases I

Operations:

- put, get, scan, delete of a row
- put, get, delete of fields of a row
- access to groups of fields of a row in a column group
- access to rows using a given timestamp



Cassandra columns

The basic unit of storage in Cassandra is the **column**

- A column is a **name-value** pair
 - The **name** behaves as the **key**
 - The **value** is a simple **byte-array**
 - A column is always stored with a **timestamp**
- The timestamp is used to expire data, resolve write conflicts, deal with stale data, and do other things

```
{  
key: "FirstName",  
value: "Alice",  
timestamp: "13645585",  
}
```



Cassandra rows

- A **row** is a collection of columns attached or linked to a key
- A collection of similar rows makes a **column family**
- Column families are defined into keyspaces
- Keyspaces have to be created so that column families can be assigned to them:

```
create keyspace game
```



Column Families

```
//column family
{
  //row
  "myAlice" : {
    firstName: "Alice",
    lastName: "Foo",
    lastGame: "2014/1/12"
  } //row
  "bobGreen" : {
    firstName: "Bob",
    lastName: "Green",
    location: "Boston"
  }
}
```



Cassandra query features

Cassandra does not have a rich query interface, because of this it is necessary to carefully choose the columns and the column families

- Cassandra client that includes the `get`, `set`, and `delete` operations
- Indexing Features
- Cassandra Query Language (CQL)



Cassandra query features

- Cassandra client

```
use games;
CREATE COLUMN FAMILY Player
WITH comparator = UTF8Type
AND key_validation_class=UTF8Type
AND column_metadata = [
    {column_name: city, validation_class: UTF8Type}
{column_name: name, validation_class: UTF8Type}];
SET Customer['myAlice']['city']='Boston';
SET Customer['myAlice']['name']='Alice Foo';
GET Customer['myAlice'];
GET Customer['myAlice']['name'];
DEL Customer['myAlice']['city'];
DEL Customer['myAlice'];
```



Cassandra Query Language

The Cassandra Query Language is a

- SQL-like query language

```
CREATE COLUMNFAMILY Player ( KEY varchar PRIMARY KEY,  
name varchar,  
city varchar);  
INSERT INTO Player (KEY,name,city)  
VALUES ('myAlice',  
       'Alice Foo',  
       'Boston');  
SELECT * FROM Customer  
SELECT name, city FROM Player  
SELECT name, city FROM Player  
WHERE city='Boston'
```



Cassandra [5]



Key-Value Databases Word





Key-Value Databases

A **simple hash table** accessible only through its **primary key**

- Basically a table with two columns: the key and the value
- The value is a blob that the data store just stores: it can be text, JSON, XML, or anything else

Operations:

- **get** the value for the key
- **put** a value for a key
- **delete** a key from the database



Key-Value Databases

- Strongly aggregate oriented
 - A lot of aggregates
 - Each aggregate has a key
- Data Model:
 - A set of (key,value) pairs
 - Value: **an aggregate instance**
- It is not possible to query/navigate the aggregate
 - the aggregate is a blob of meaningless bits
- Access to an aggregate
 - lookup based on **its key**



Example

[A first data representation in Oracle NoSQL]

key	value
/Player/mary1994/-/username	mary1994
/Player/mary1994/-/firstName	Mary
/Player/mary1994/-/lastName	Wilson
/Player/mary1994/-/games	...

[A second data representation in Oracle NoSQL]

key	value
/Player/mary1994/-/username	mary1994
/Player/mary1994/-/firstName	Mary
/Player/mary1994/-/lastName	Wilson
/Player/mary1994/-/games/0/id	2345
/Player/mary1994/-/games/0/oppont	rick_the_good
/Player/mary1994/-/games/0/gameDetails	...
/Player/mary1994/-/games/1/id	7425
/Player/mary1994/-/games/1/oppont	ann_x
/Player/mary1994/-/games/1/gameDetails	...



Example

[Key-value per object in Redis]

key	value
Player:mary1994	"username":"mary1994", "firstName":"Mary", ...

[Key-value per field in Redis]

key	value
Player:mary1994/username	mary1994
Player:mary1994/firstName	Mary
Player:mary1994/lastName	Wilson
Player:mary1994/games	...

[Key-hash per field in Redis]

key	value
4*Player:mary1994	username:mary1994 firstName:Mary lastName:Wilson games:....



Query Features

General

- Only query by the **key**
- It is not possible to use some attribute of the value column
- If the key is not known:
 - Some systems allow the search inside the value
 - e.g., Riak Search
 - The key needs to be suitably chosen
 - e.g., session ID for storing session data



Example

[Key-value per object in Redis]

key	value
Player:mary1994	"username":"mary1994", "firstName":"Mary", ...

[Key-value per field in Redis]

key	value
Player:mary1994/username	mary1994
Player:mary1994/firstName	Mary
Player:mary1994/lastName	Wilson
Player:mary1994/games	...

[Key-hash per field in Redis]

key	value
4*Player:mary1994	username:mary1994 firstName:Mary lastName:Wilson games:....



Query Features

General

- Only query by the **key**
- It is not possible to use some attribute of the value
- If the key is not known:
 - Some systems allow the search inside the value
 - e.g., Riak Search
 - The key needs to be suitably chosen
 - e.g., session ID for storing session data



Redis

Database Structure

Redis is an open source in-memory data structure store

It can be used as a:

- database
- cache
- message broker



Redis

Database Structure

Redis is an open source in-memory data structure store

It can be used as a:

-
-
- message broker



Redis

Databases Structure

- Stores key-value pairs

The value is not limited to a simple string:

- Binary-safe strings
- Lists
- Sets
- Sorted sets
- Hashes
- Bit arrays
- HyperLogLogs



Redis

Databases Structure

- Stores key-value pairs

The value is not limited to a simple string:

- Binary-safe strings



Redis

Databases Structure

- Stores key-value pairs

The value is not limited to a simple string:

- Binary-safe strings
- Lists
- Sets
- Sorted sets
- Hashes
- Bit arrays
- HyperLogLogs

Redis

Databases Structure

- Stores key-value pairs

The value is not limited to a simple string:

- HyperLogLogs



Redis

HyperLogLog

probabilistic data structure used to count unique values

Counting unique values with exact precision requires an amount of memory proportional to the number of unique values:

- trades memory consumption for precision



Query Features

Redis commands

provides basic **store, fetch, delete** operations

Redis APIs

There are modules/APIs that allow to integrate the connection to Redis into the code



Query Features

Redis commands

`redis-cli` is the Redis command line interface

- send commands to Redis
- read the replies sent by the server
- operations run directly from the terminal



Query Features

Simple operations

```
SET "my key" "my string to store"  
GET "my key"
```



Key-Value Databases Word

```
[funny-ecran:src francy$ redis-cli
[127.0.0.1:6379> SET a b
OK
[127.0.0.1:6379> get a
"b"
127.0.0.1:6379> ]
```



Query Features

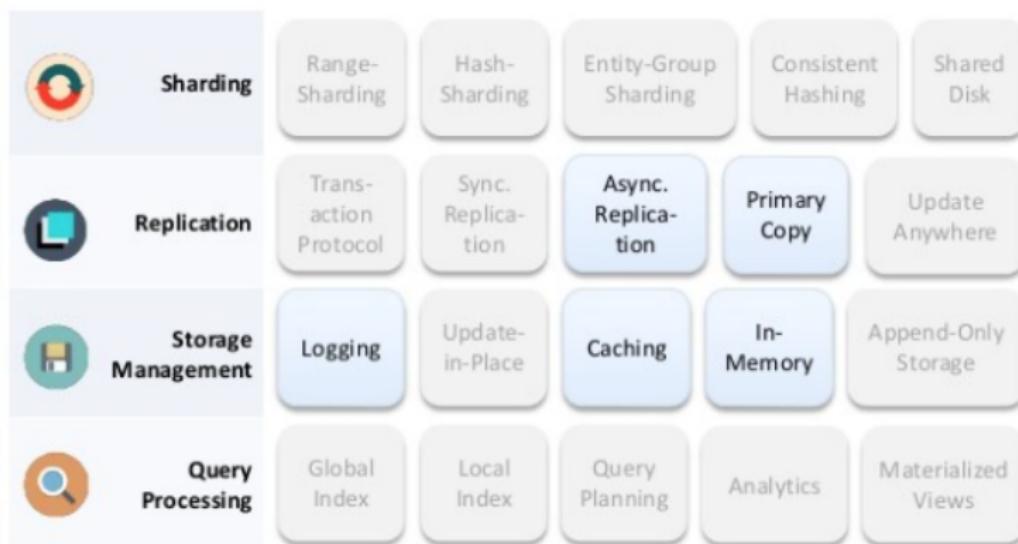
Other commands

- GETSET "key" "value"
- RPUSH "key" "value" ["value" ...]
- ...

<https://redis.io/commands>



Redis [5]



Graph Database Word





Graph Databases

- Different point of view
 - Complex relations require complex joins
- Goal:
 - Capture data consisting of complex relations
 - Data naturally modeled as graphs
 - Social Networks, Web data, etc.
 - No joins but navigations



Graph Databases

Allow to store entities and relationships between entities

- Nodes
- Edges

A node is an instance of an entity and the relations between entities are modeled as edges

- Nodes and edges can have **properties** and **types**
- Edges have **directional significance**

Operations:

- **Traversals**
 - edges following on the basis of some condition
- **Shortest path**
 - going from one node to another following the shortest path
- **Optimal path**
 - we want to include some nodes

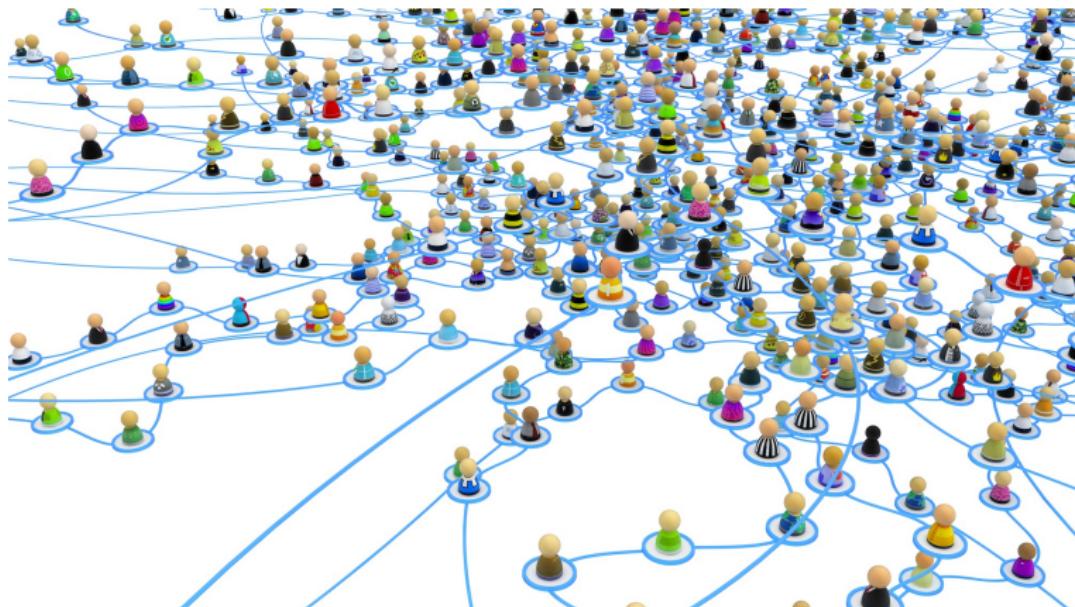


Graph Databases

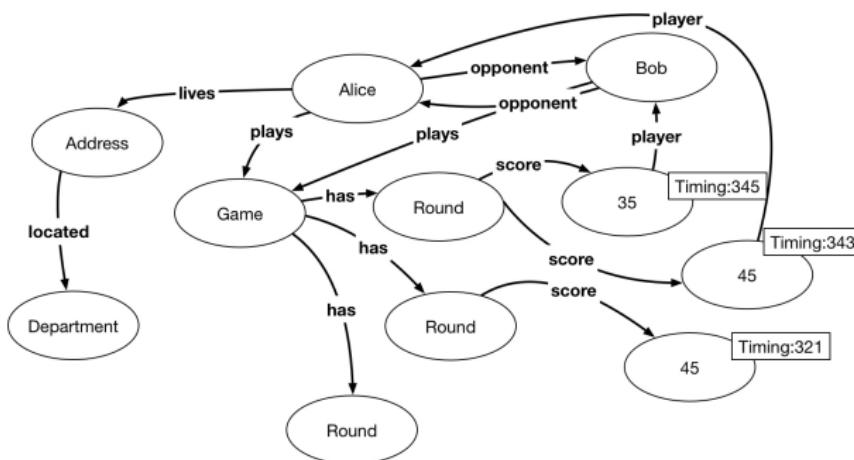
- Not aggregate oriented
 - data represented as nodes connected with edges
- Data Model:
 - Nodes connected by edges
 - Many variations:
 - Neo4j stores java objects to nodes and edges in a schemaless fashion
 - Infinite Graph stores Java objects, which are subclasses of built in types as nodes and edges
 - FlockDB simply nodes and edges with no additional attributes
- Access to data:
 - Navigation through the network of edges
 - Need of a starting node
 - Nodes can be indexed by an attribute



Graph Databases



Example





Neo4j data model

Neo4j is a highly scalable native graph database

- Represents data in nodes, relationships and properties
- Both nodes and relationships contain properties
- Relationships connect nodes
- Properties are key-value pairs
- Relationships have directions: unidirectional and bidirectional



Neo4j data model

- How we can create the nodes and assign them properties:

```
Node alice = graphDb.createNode();
alice.setProperty("name", "Alice");
```

```
Node bob = graphDb.createNode();
bob.setProperty("name", "Bob");
```



Relationships

- Relationships are first class citizens in graph databases:
 - Type
 - Start node
 - End node
 - Properties
- The properties can be used to query the graph
- Design work to model relationships



Neo4j data model

- How we can create relationships:

```
alice.createRelationshipTo(bob, OPPONENT);
```

```
bob.createRelationshipTo(alice, OPPONENT);
```



Query features

General

- Specific graph query languages
 - **Gremlin**: a domain specific language for traversing graphs
 - **Cypher** an SQL-like query language used in Neo4j

Neo4j

- Cypher
- Features for:
 - traversing the graph
 - querying the graph properties
 - navigating nodes using language bindings



Traversals

General

- Powerful when it is necessary to traverse the graph at any depth and specify a starting node for the traversal
- Very powerful when you do not know how manage the depth of the query
- Very powerful when you have to go more levels down
- It is also possible to make the traversal top-down



Traversals

Neo4j

- The `friendsTraverser`

- finds all the nodes that are related to Alice where the relationship type is `OPPONENT`
- The nodes can be at any depth—opponent

```
Node alice = nodeIndex.get("name", "Alice").getSingle();
Traverser friendsTraverser = alice.traverse(Order.BREADTH_FIRST,
StopEvaluator.END_OF_GRAPH,
ReturnableEvaluator.ALL_BUT_START_NODE,
EdgeType.OPPONENT,
Direction.OUTGOING);
```



Traversals

Neo4j

Suppose that we want to discover if Alice and Bob have ever been opponents in a game

- We retrieve **all** the paths (all the times they were opponents)

```
Node alice = nodeIndex.get("name", "Alice").getSingle();
Node bob = nodeIndex.get("name", "Bob").getSingle();
PathFinder<Path> finder = GraphAlgoFactory.allPaths(
    Traversal.expanderForTypes(FRIEND, Direction.OUTGOING)
    ,MAX_DEPTH);
Iterable<Path> paths = finder.findAllPaths(alice, bob);
```

Finds all the paths and their distance



Traversals

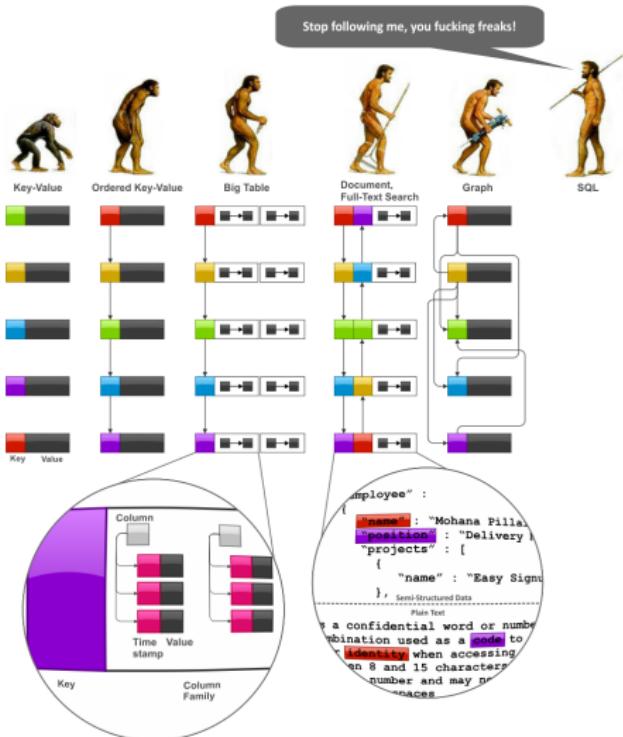
- If there are multiple paths the **shortest** path is found:

```
PathFinder<Path> finder = GraphAlgoFactory.shortestPath(  
    Traversal.expanderForTypes(FRIEND, Direction.OUTGOING)  
        , MAX_DEPTH);  
Iterable<Path> paths = finder.findAllPaths(barbara, jill);
```

Finds the shortest path



Comparing [7]...





References |

-  Francesca Bugiotti, Luca Cabibbo, Paolo Atzeni, and Riccardo Torlone.
Database design for nosql systems.
In *Conceptual Modeling - 33rd International Conference, ER 2014, Atlanta, GA, USA, October 27-29, 2014. Proceedings*, pages 223–231, 2014.
-  Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Michael Burrows, Tushar Chandra, Andrew Fikes, and Robert Gruber.
Bigtable: A distributed storage system for structured data (awarded best paper!).
In *7th Symposium on Operating Systems Design and Implementation (OSDI '06), November 6-8, Seattle, WA, USA*, pages 205–218, 2006.



References II

-  E. F. Codd.
Relational completeness of data base sublanguages.
IBM Research Report, RJ987, 1972.
-  Eric Evans.
Domain-Driven Design: Tacking Complexity In the Heart of Software.
Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2003.
-  Felix Gessert and Norbert Ritter.
Scalable data management: Nosql data stores in research and practice.
In *2016 IEEE 32nd International Conference on Data Engineering (ICDE)*, pages 1420–1423, 2016.



References III

-  Katarina Grolinger, Wilson Higashino, Abhinav Tiwari, and Miriam Capretz.
Data management in cloud environments: Nosql and newsql data stores.
Journal of Cloud Computing: Advances, Systems and Application, 2, 12 2013.
-  NoSQL Data Modeling Techniques.
<https://highlyscalable.wordpress.com/2012/03/01/nosql-data-modeling-techniques/>.
Accessed 2016.